

CharacterBERT and Self-Teaching for Improving the Robustness of Dense Retrievers on Queries with Typos


Shengyao Zhuang & Guido Zuccon

{s.zhuang,g.zuccon}@uq.edu.au

ielab, The University of Queensland, Australia

www.ielab.io

BERT-based rankers cannot deal with queries with typos

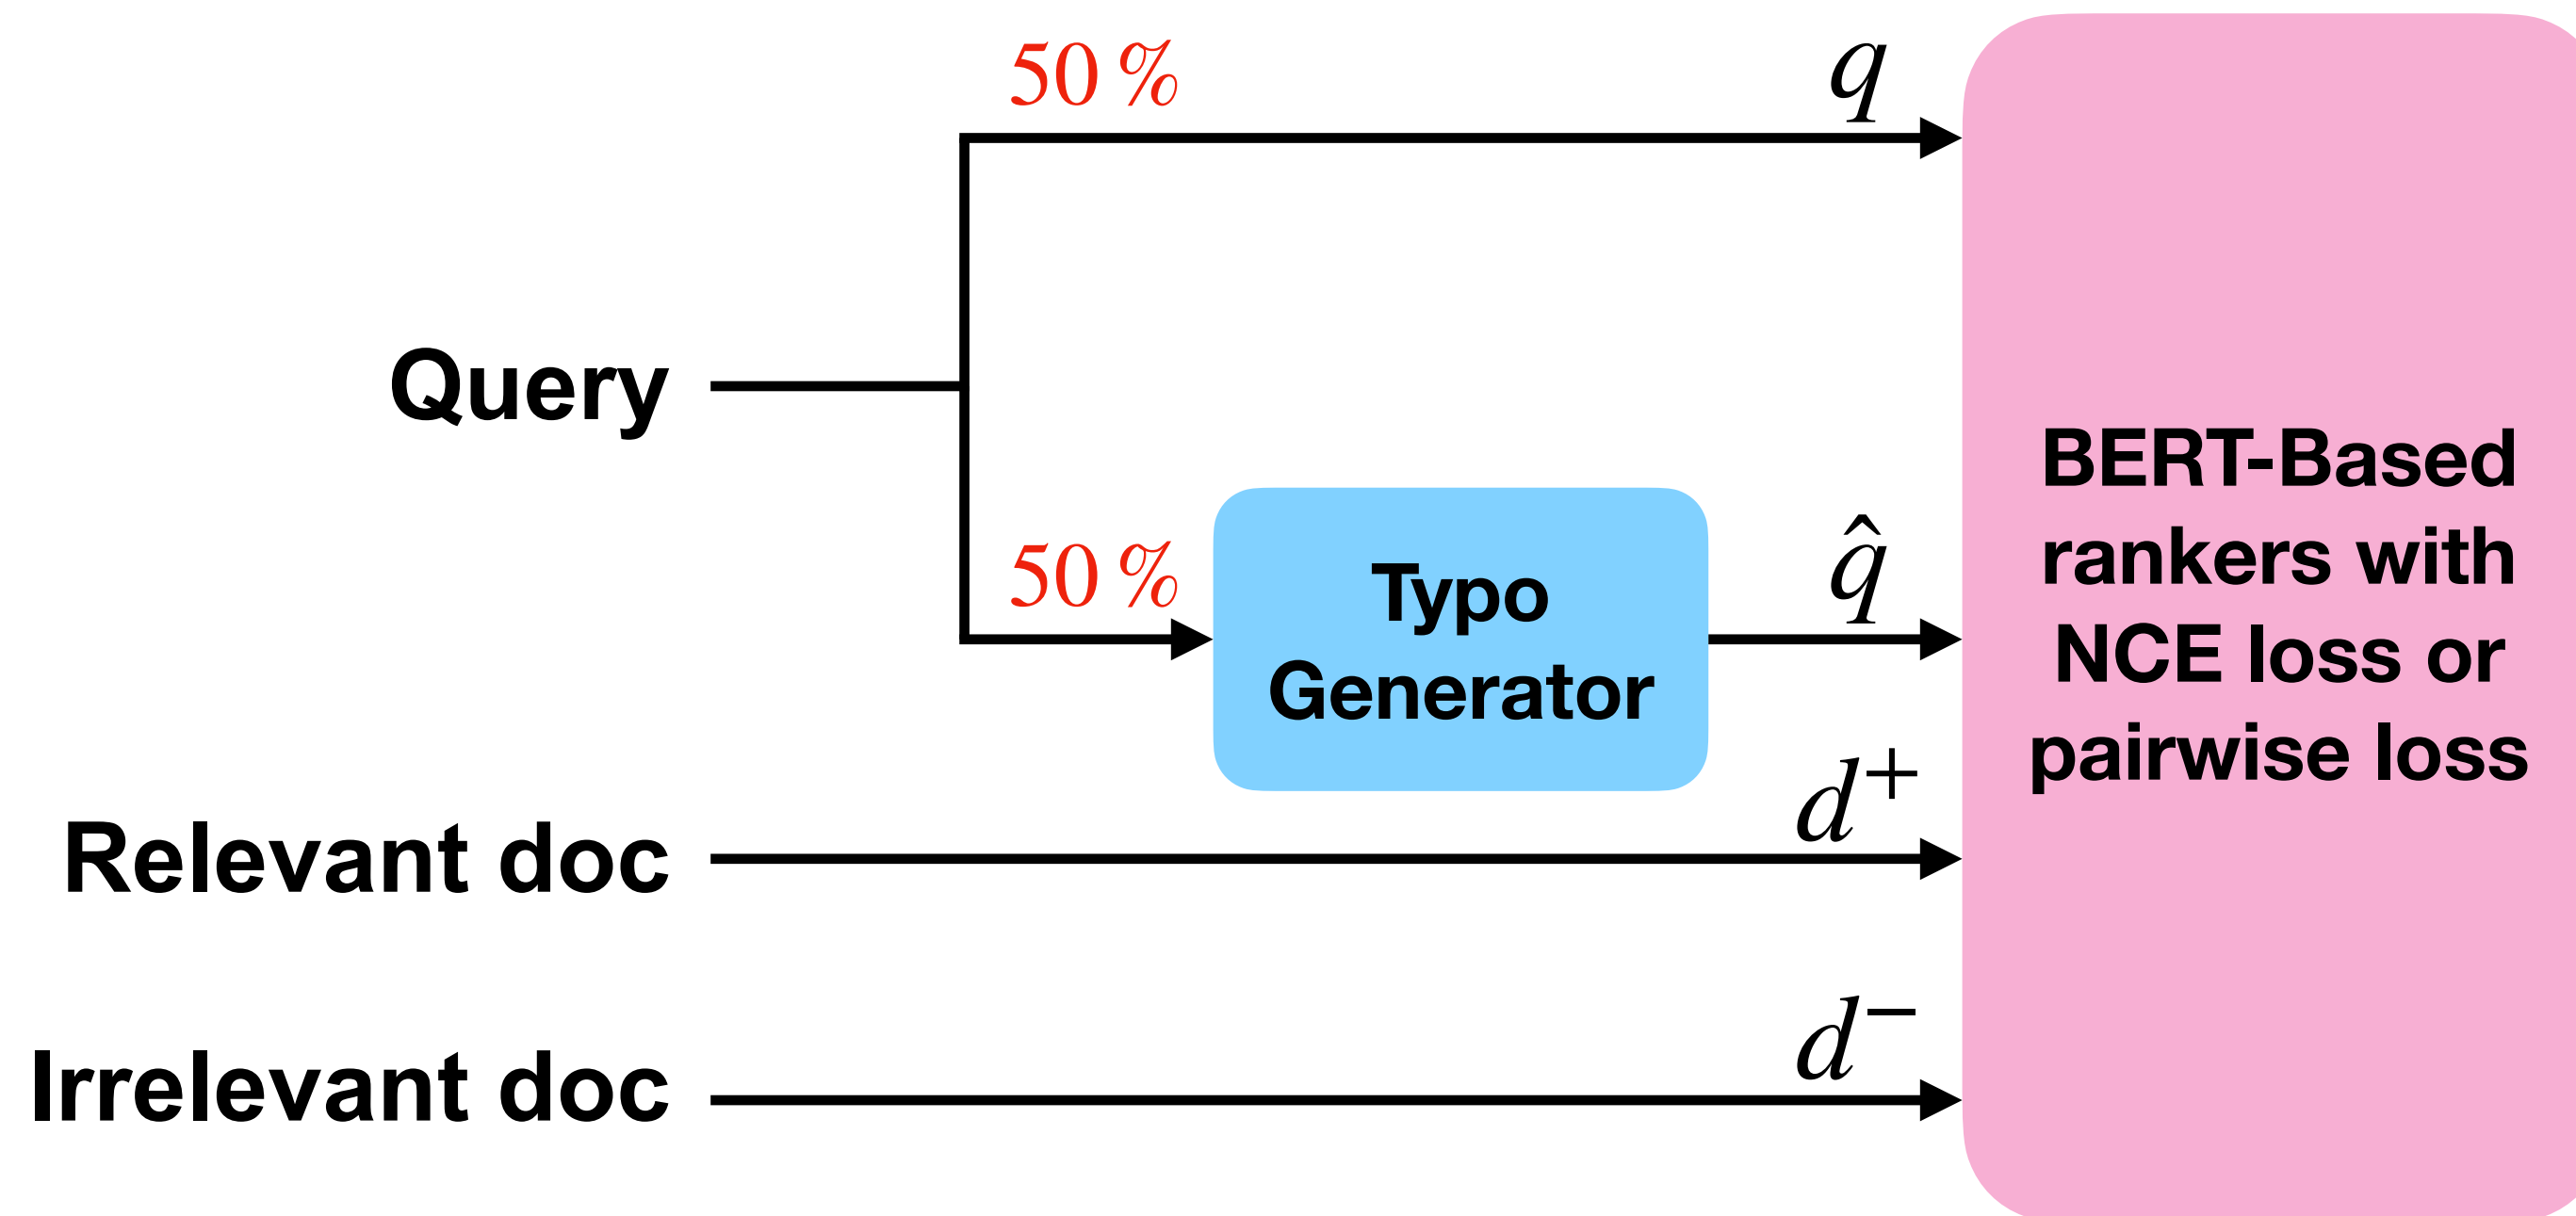


Both DR and monoBERT perform poorly on queries with typos




Zhuang, Zuccon, "Dealing with Typos for BERT-based Passage Retrieval and Ranking", EMNLP 2021

Prior work: Typos-aware training

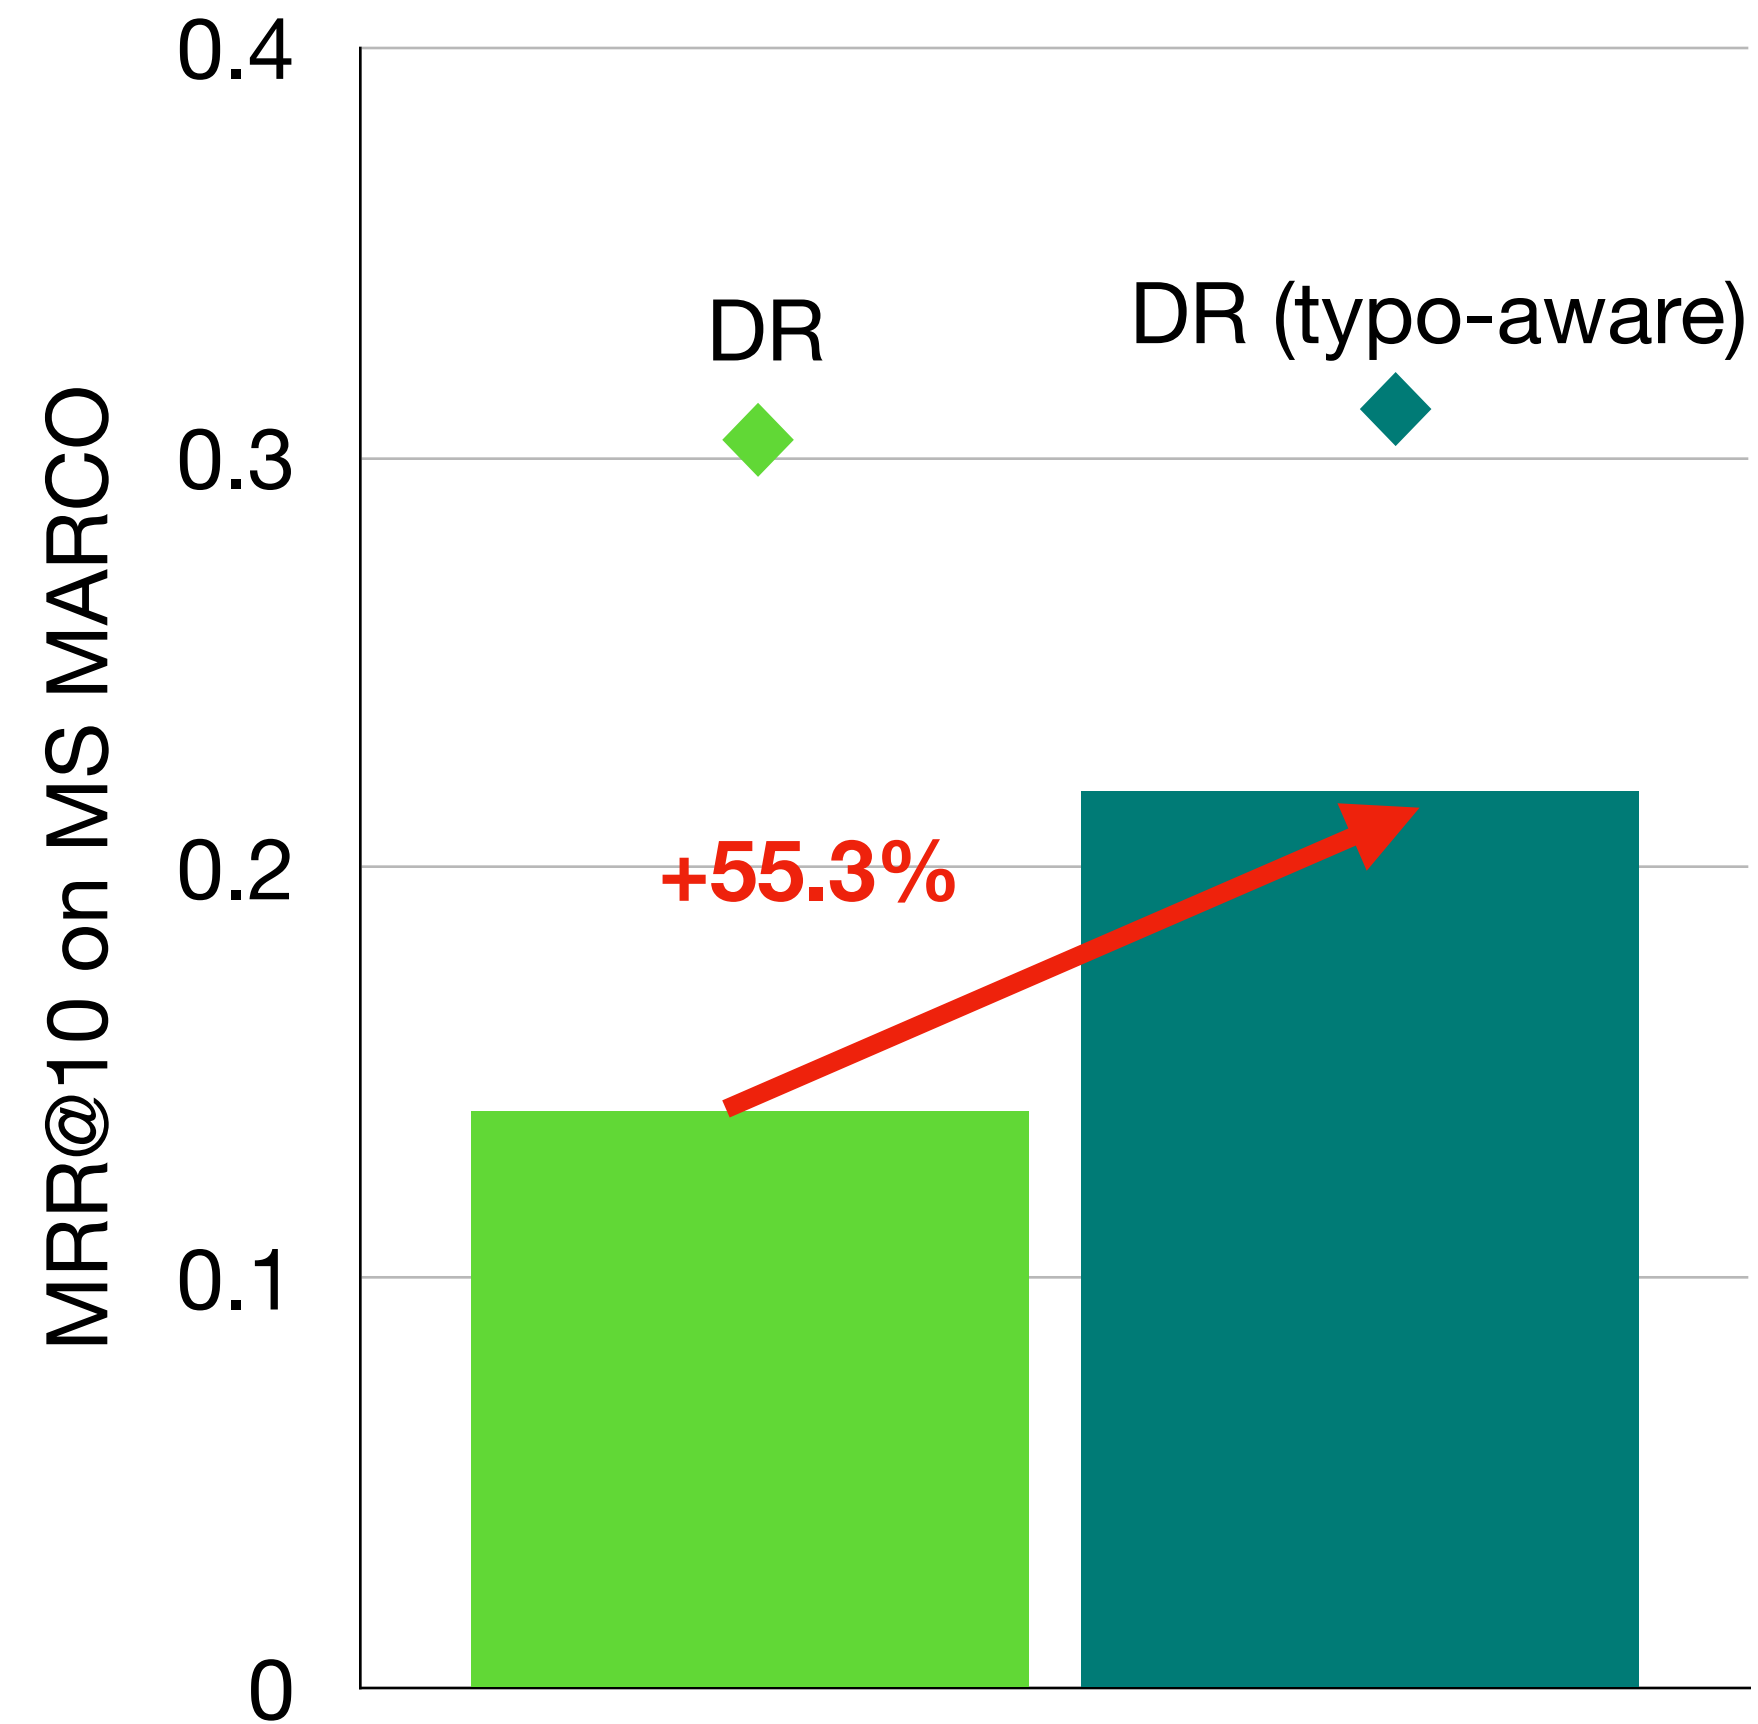


Data Augmentation strategy: inject typos in training queries




Zhuang, Zuccon, "Dealing with Typos for BERT-based Passage Retrieval and Ranking", EMNLP 2021

Prior work: Typos-aware training



- Data Augmentation improves effectiveness
- Typo-aware training as effective as traditional training on queries without typos
- Typo-aware training more effective on queries with typos

In this paper...



- Focus on **Dense Retrievers**
- **Why** Dense Retrievers cannot deal with typos in queries?
- **How** can we **improve** Dense Retriever's robustness to typos in queries?

Why BERT-based DRs cannot deal with queries with typos?

- BERT is pre-trained on curated text, and MS MARCO dataset has no or very few queries with typos.
- WordPiece Tokenization based on small vocabulary which contains common terms + common subtokens
- What is the difference in output from the WordPiece Tokenization in presence of a typo?

'information retrieval' $\xrightarrow{\text{tokenize}}$ ['information', 'retrieval'] $\xrightarrow{\text{input_ids}}$ [2592, 26384]


'inf**ro**mation retrieval' $\xrightarrow{\text{tokenize}}$ ['in', 'fr', 'oma', 'tion', 'retrieval'] $\xrightarrow{\text{input_ids}}$ [1999, 19699, 9626, 3508, 26384]

Why BERT-based DRs cannot deal with queries with typos?

- BERT is pre-trained on curated text, and MS MARCO dataset has no or very few queries with typos.
- WordPiece Tokenization based on small vocabulary which contains common terms + common subtokens
- What is the difference in output from the WordPiece Tokenization in presence of a typo?

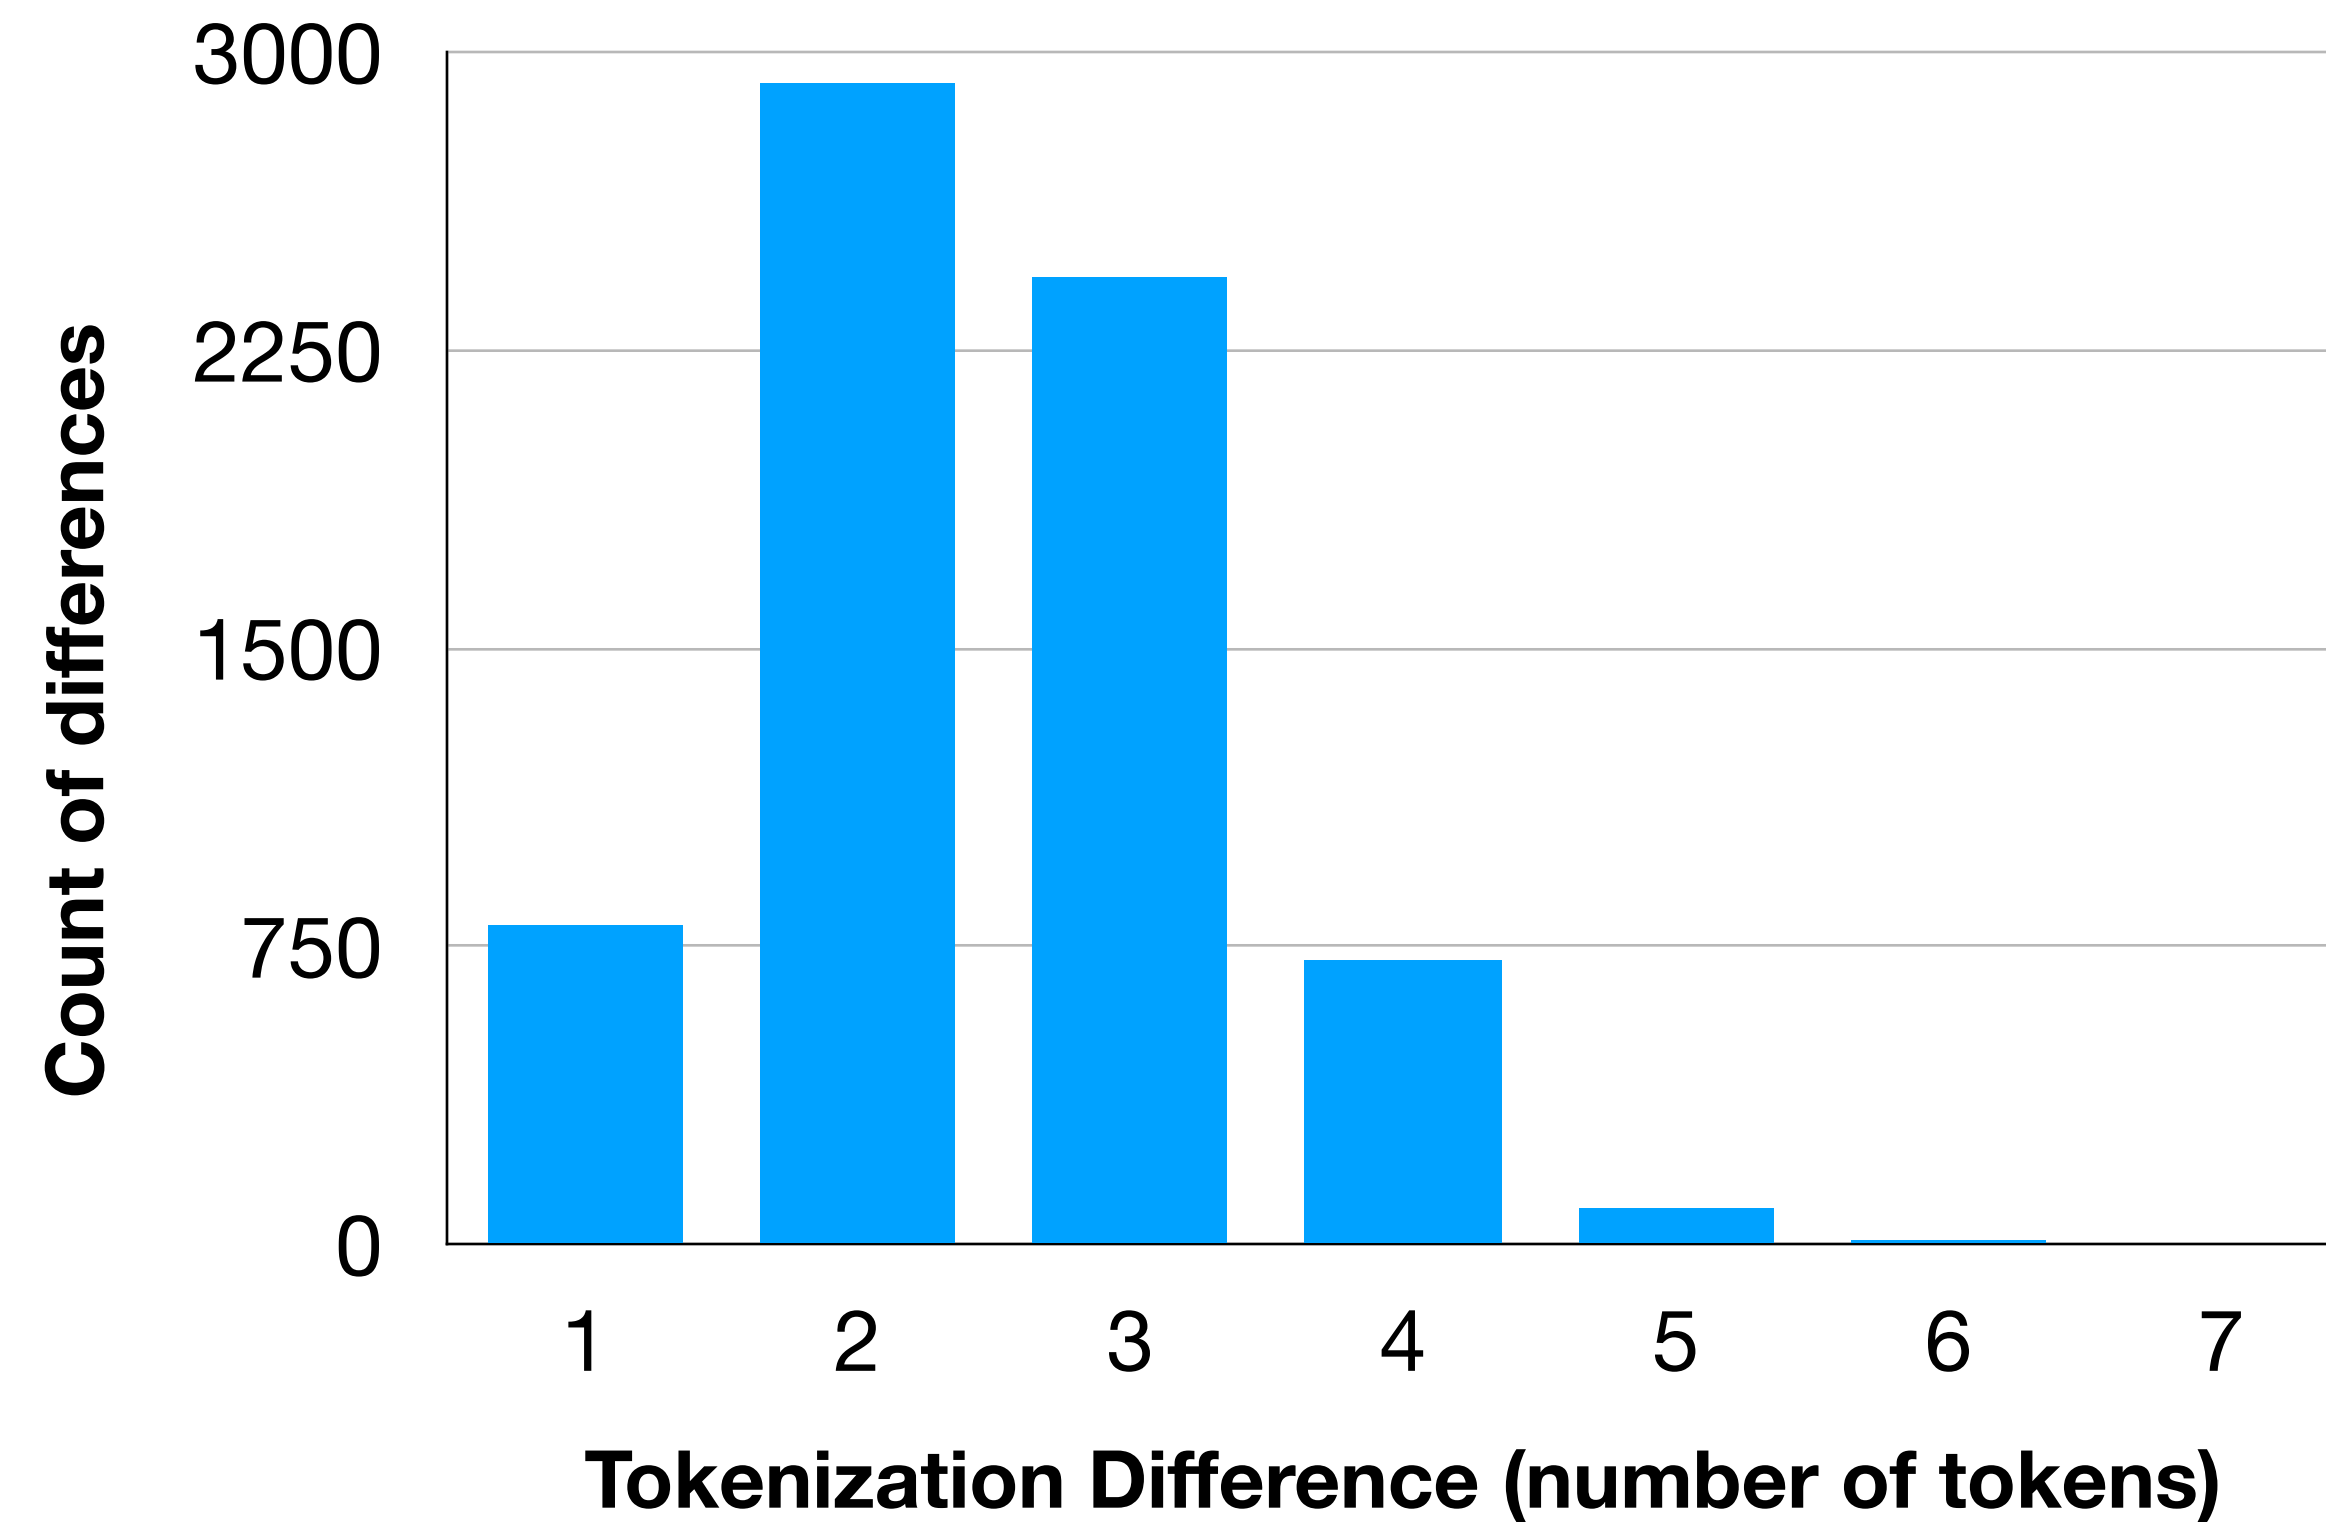
'information retrieval' $\xrightarrow{\text{tokenize}}$ ['information', 'retrieval'] $\xrightarrow{\text{input_ids}}$ [2592, 26384]

'inf~~ro~~mation retrieval' $\xrightarrow{\text{tokenize}}$ ['in', 'fr', 'oma', 'tion', 'retrieval'] $\xrightarrow{\text{input_ids}}$ [1999, 19699, 9626, 3508, 26384]




A typo resulted in the query being represented by 4 additional tokens (and lost 1)

How large are the tokenization differences caused by typos in queries?

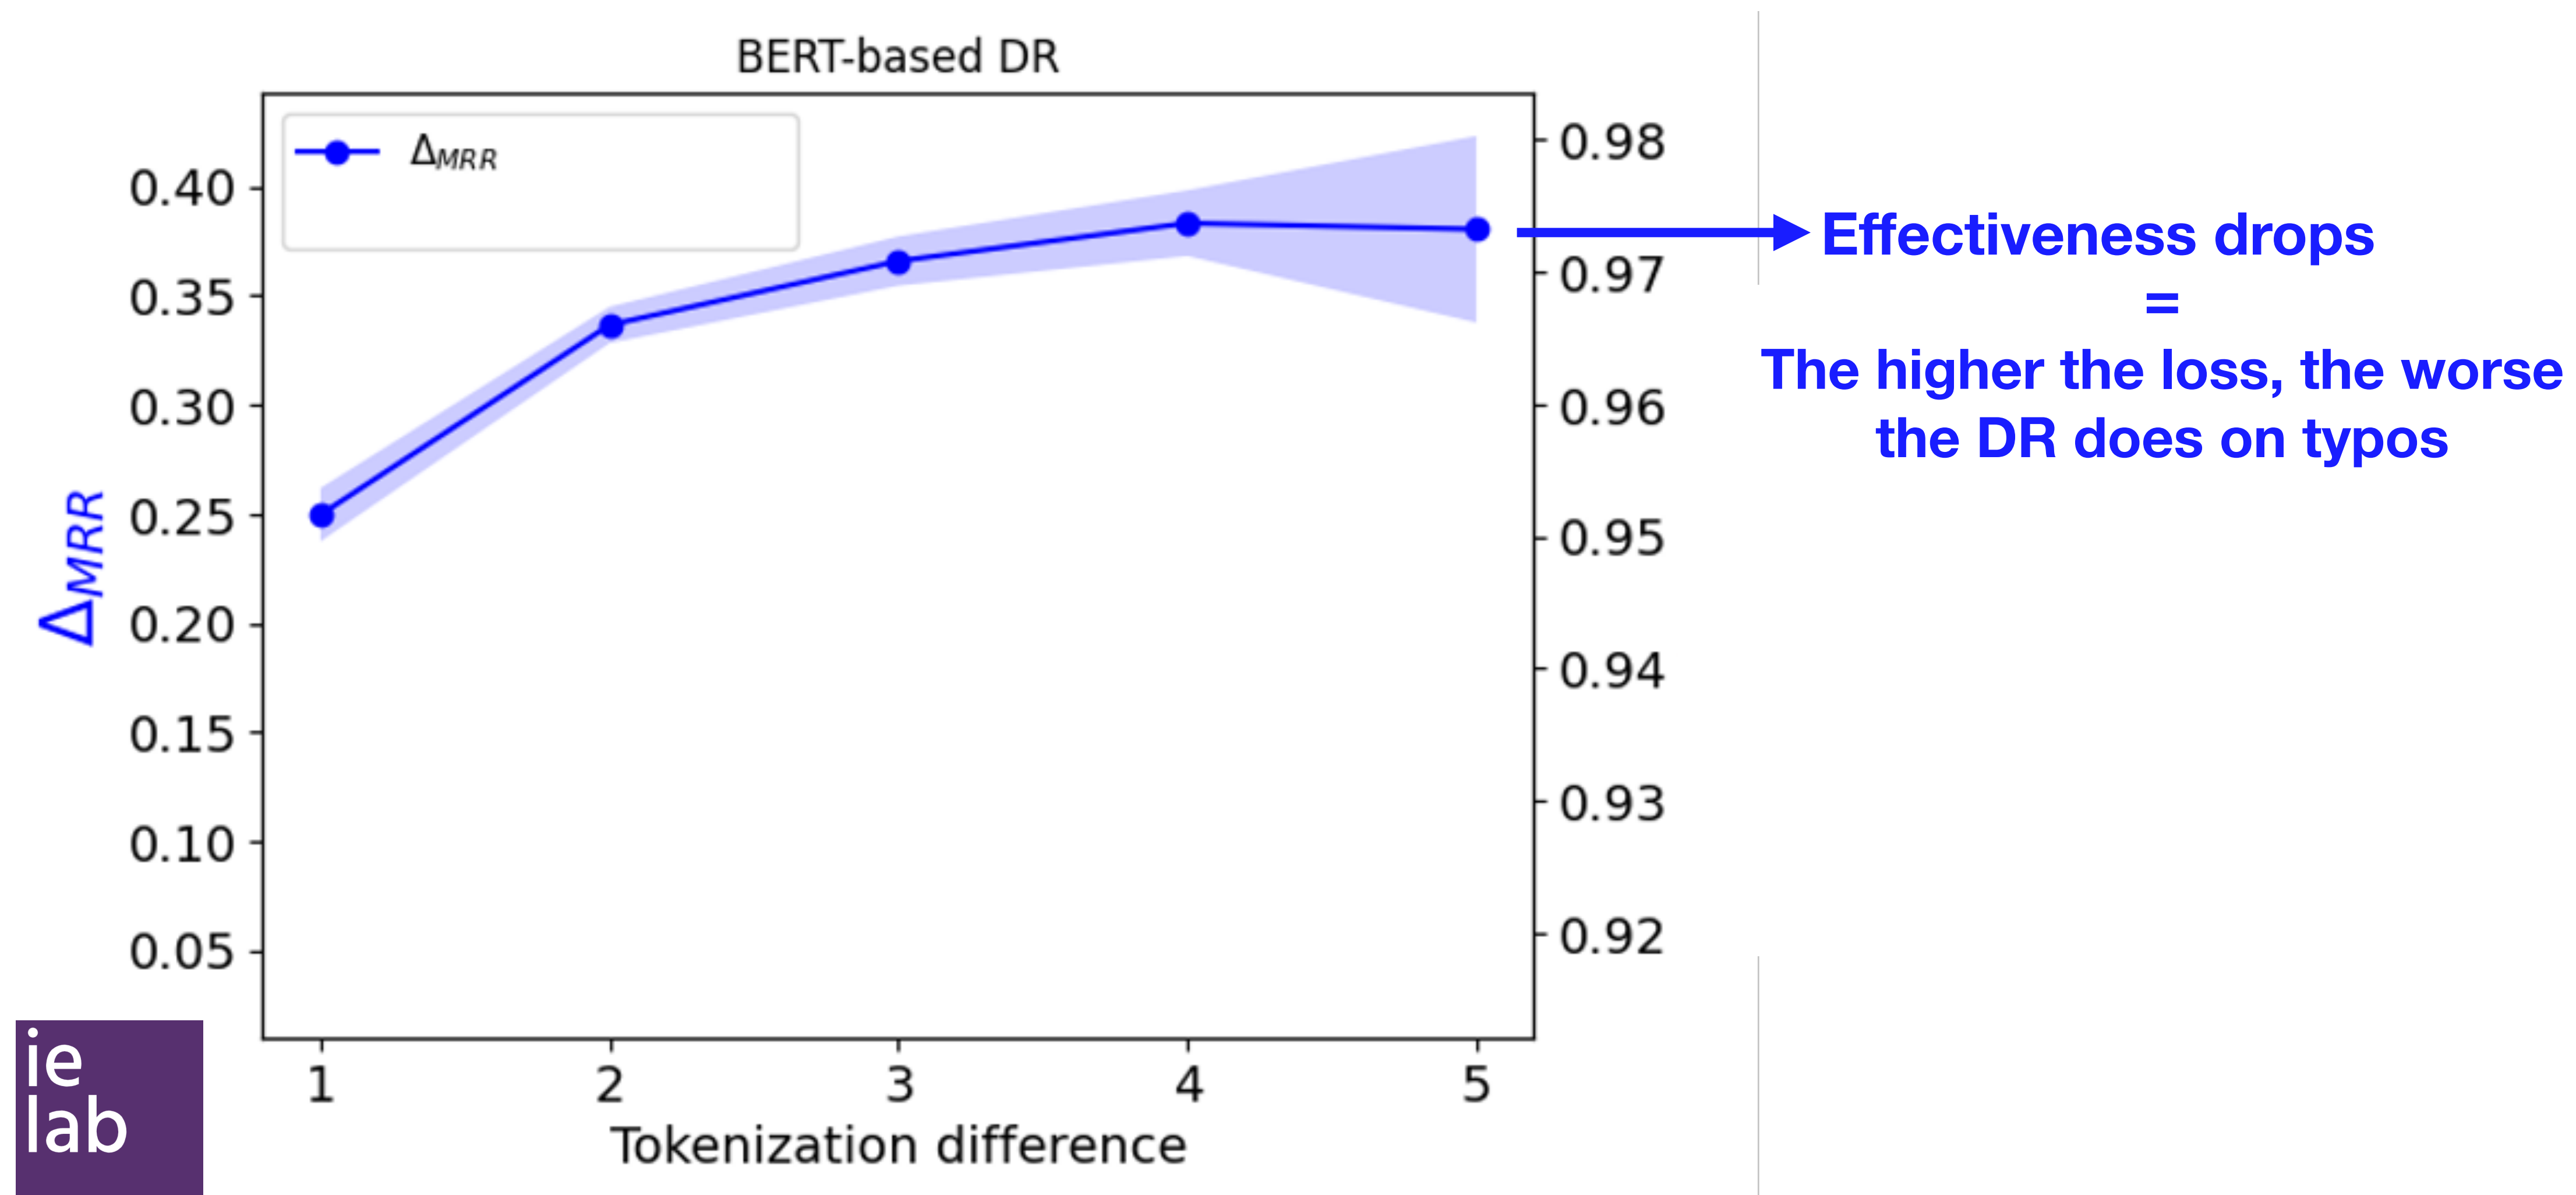


How large are the tokenization differences caused by typos in queries?

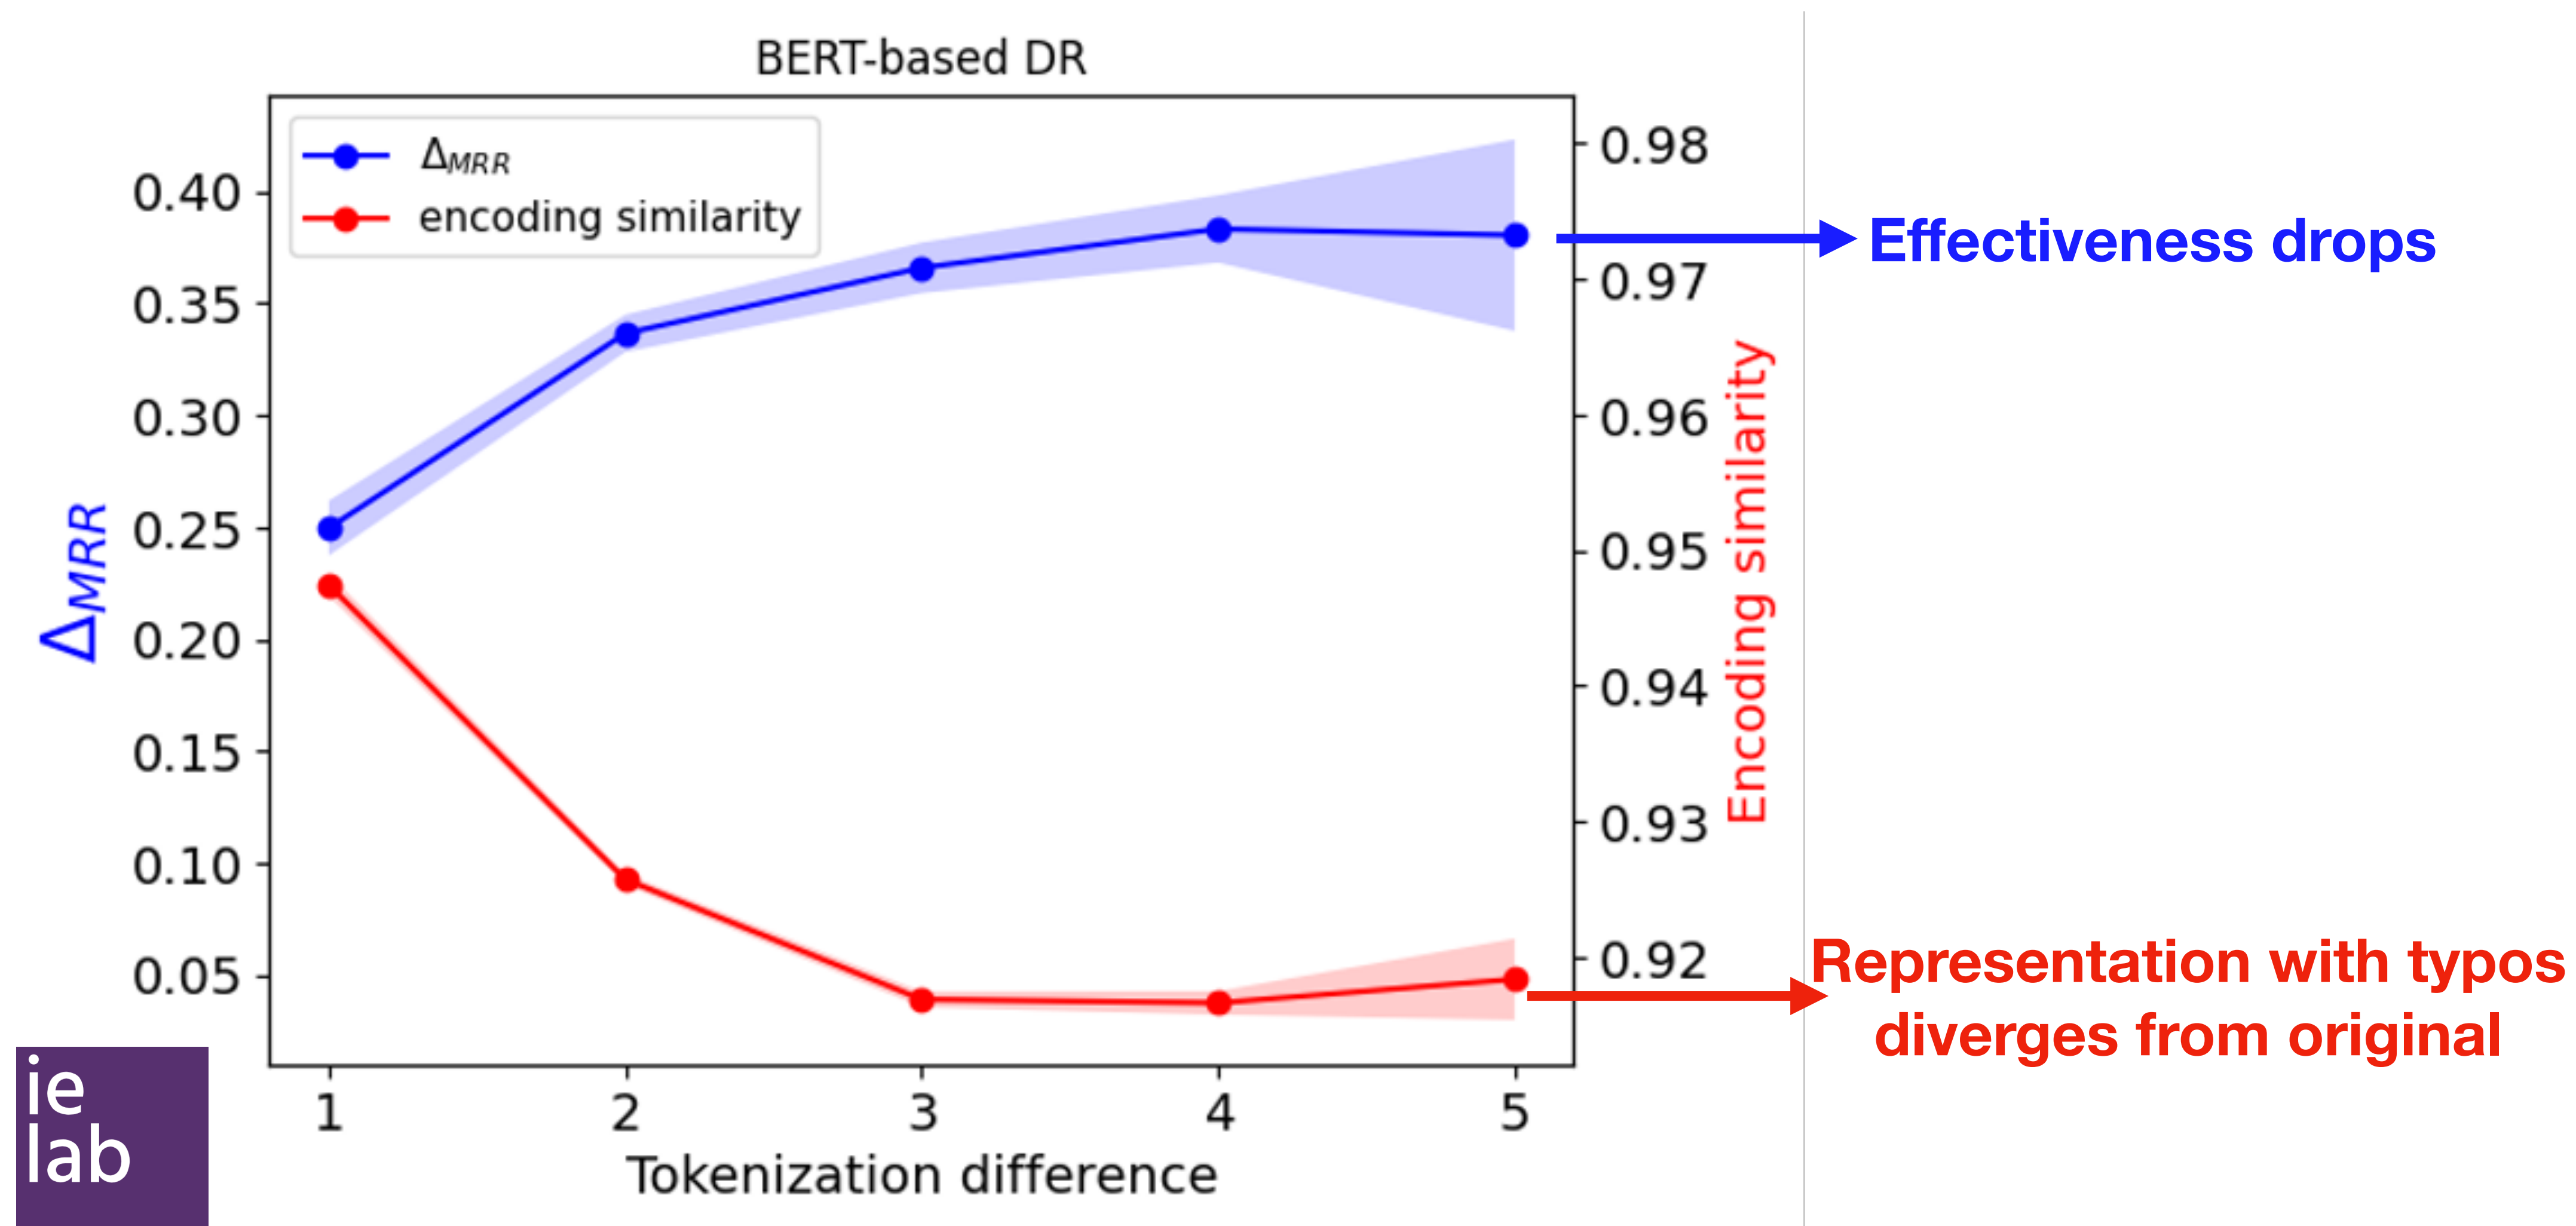


What are the effects of such tokenization difference?

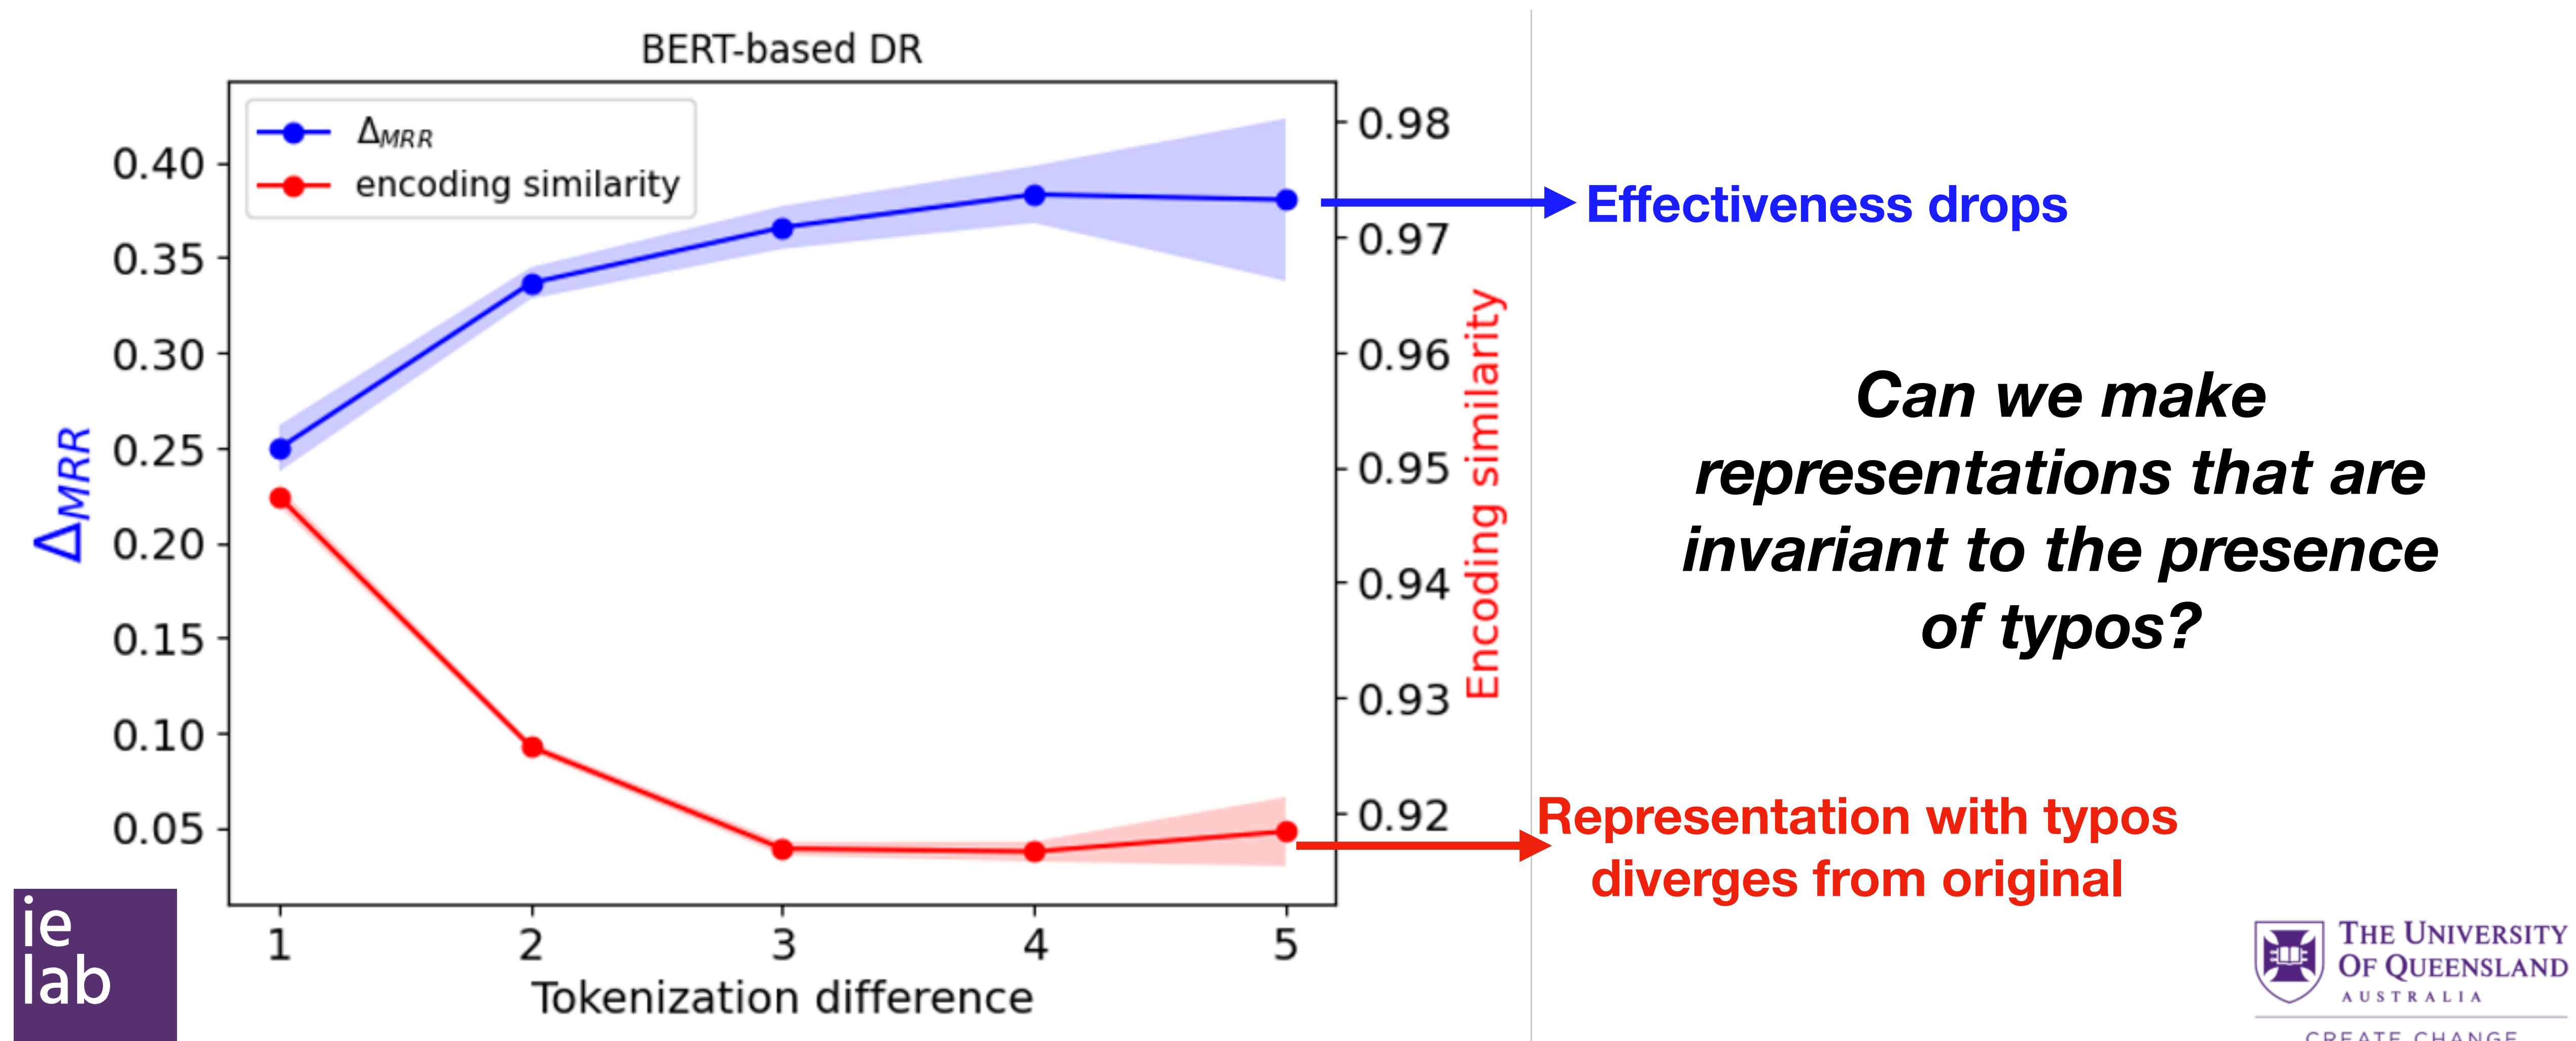
Tokenization Difference produces effectiveness losses



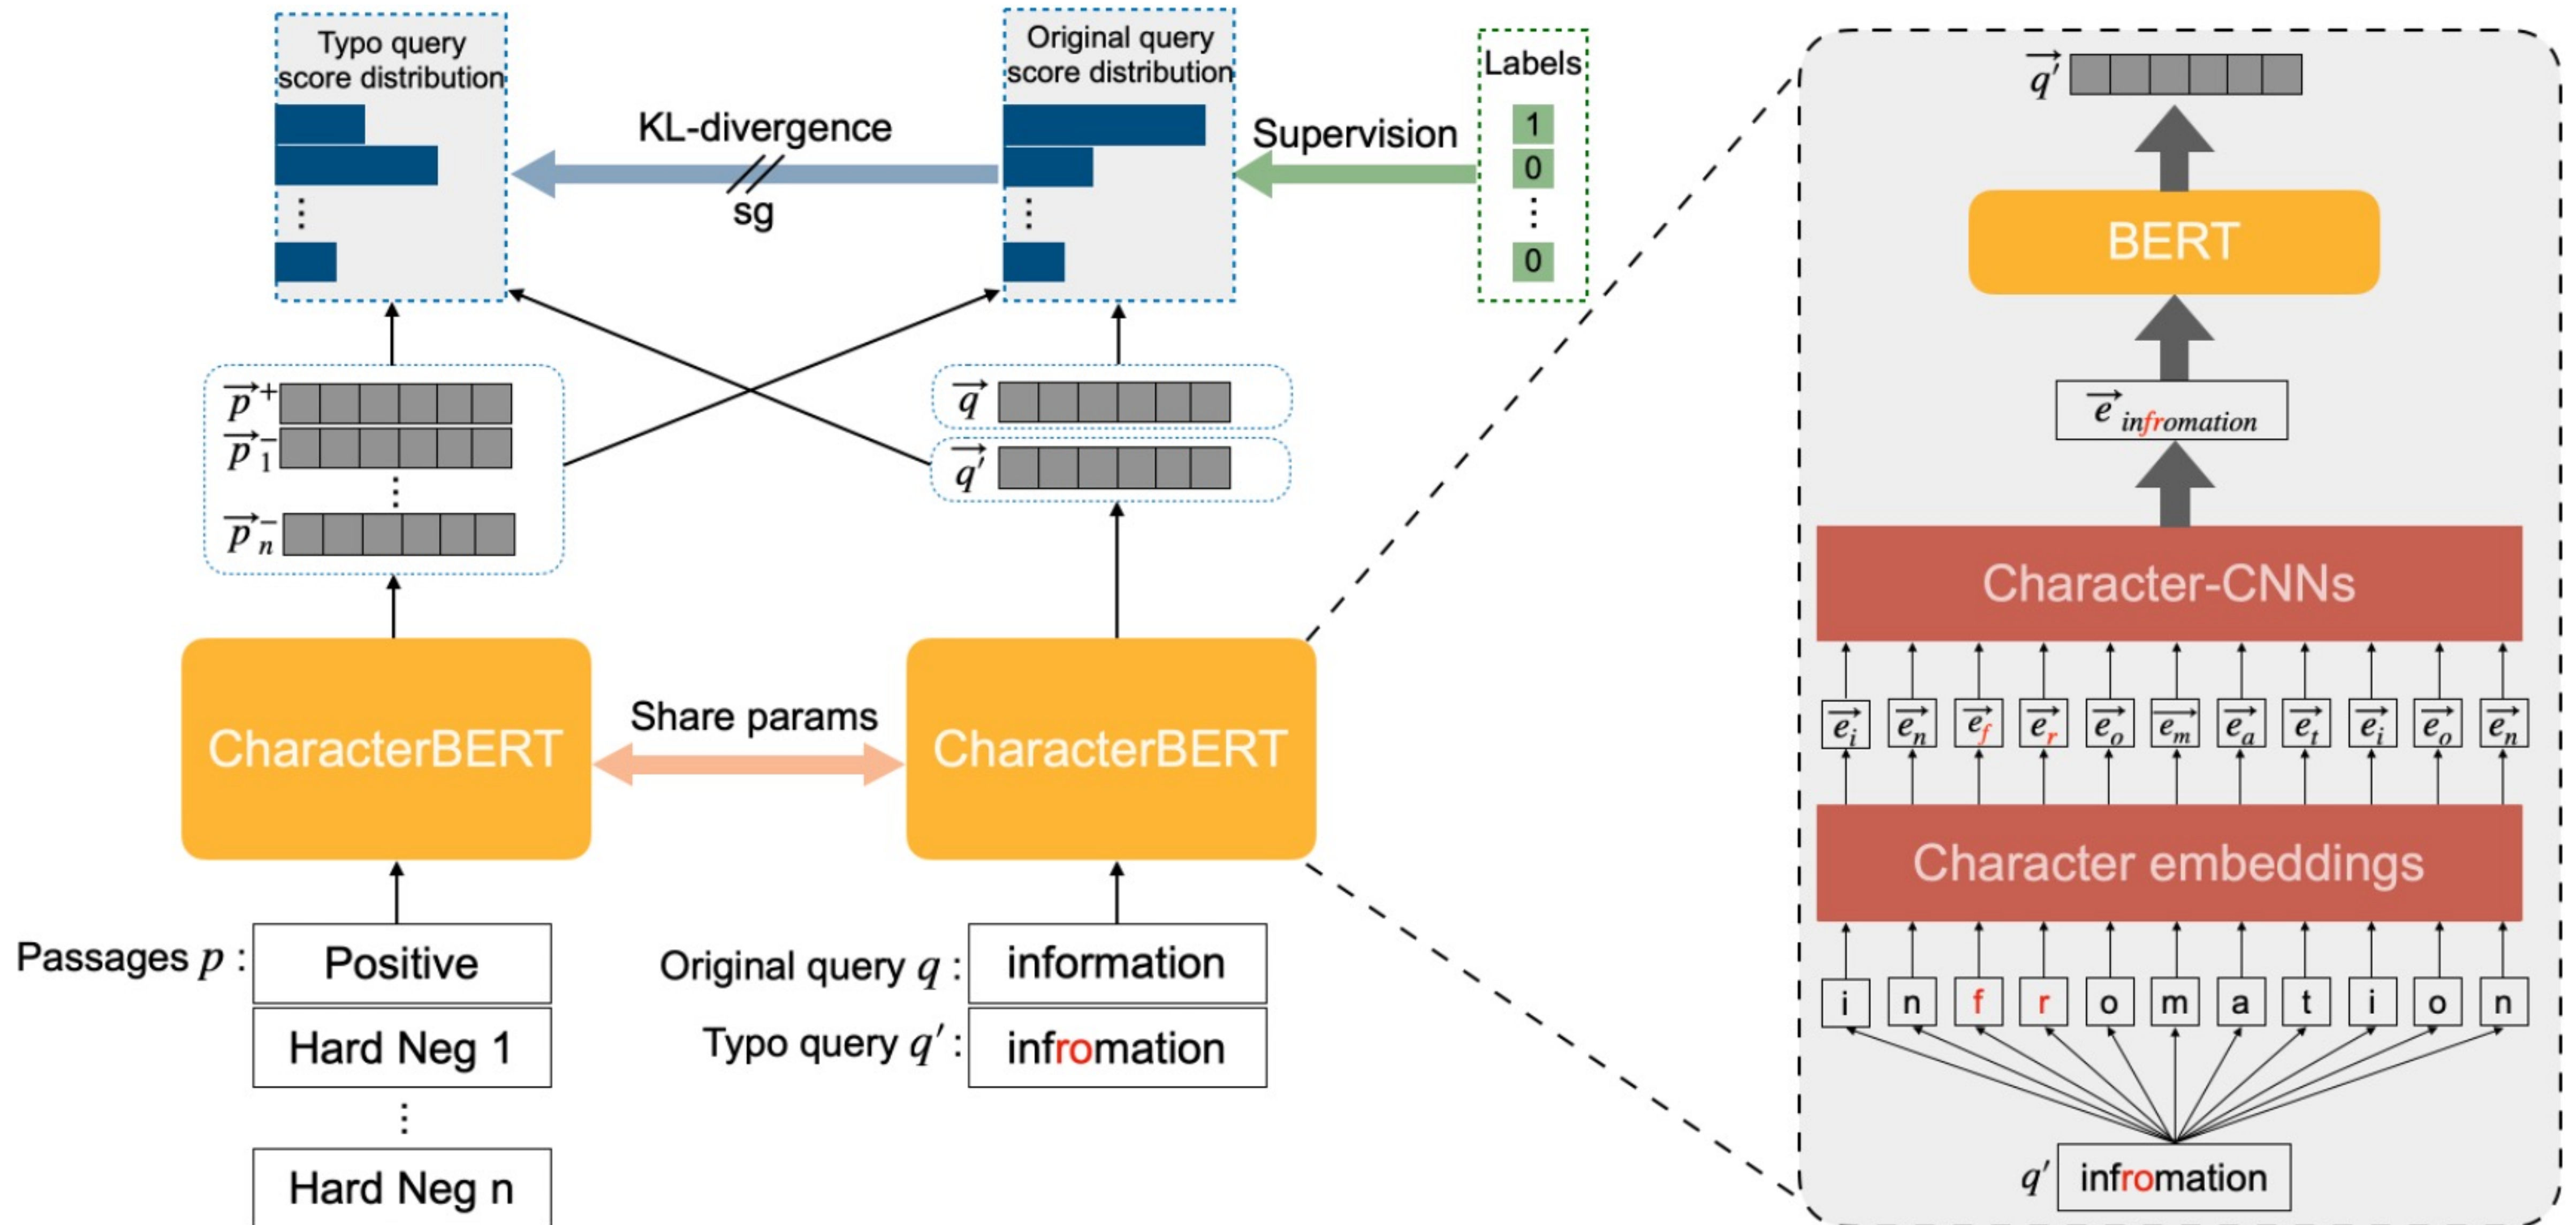
Tokenization Difference produces effectiveness losses & encodings different from query without typo



Tokenization Difference produces effectiveness losses & encodings different from query without typo



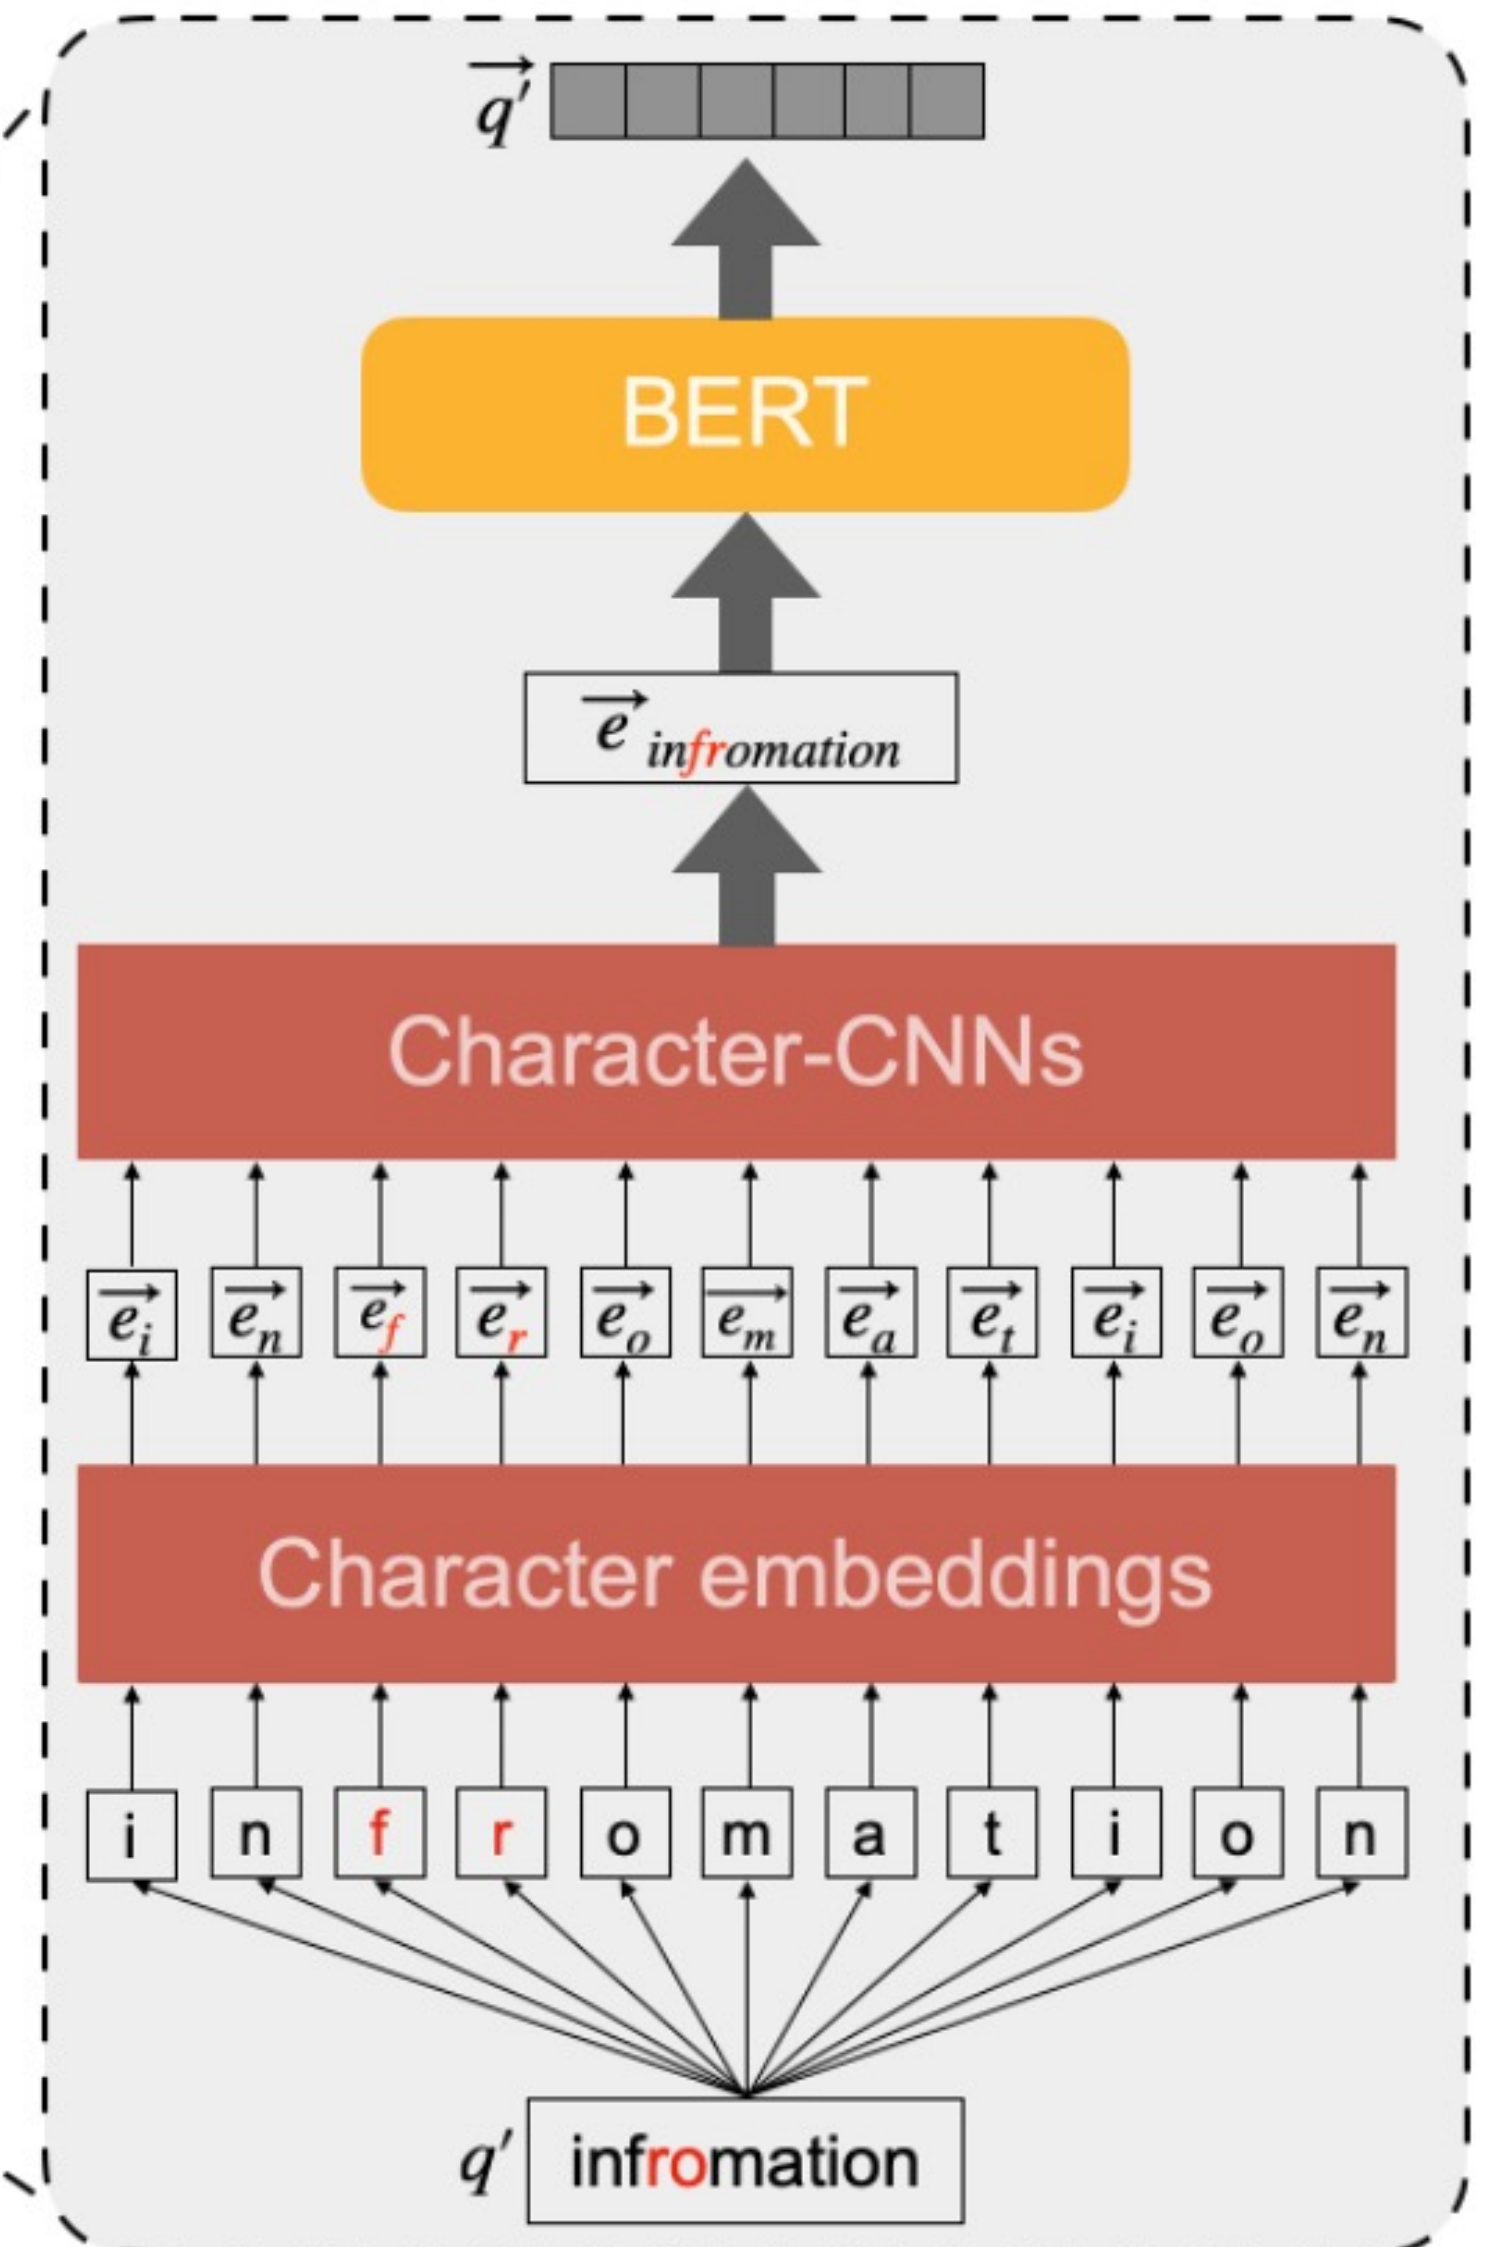
CharacterBERT-DR + Self-Teaching



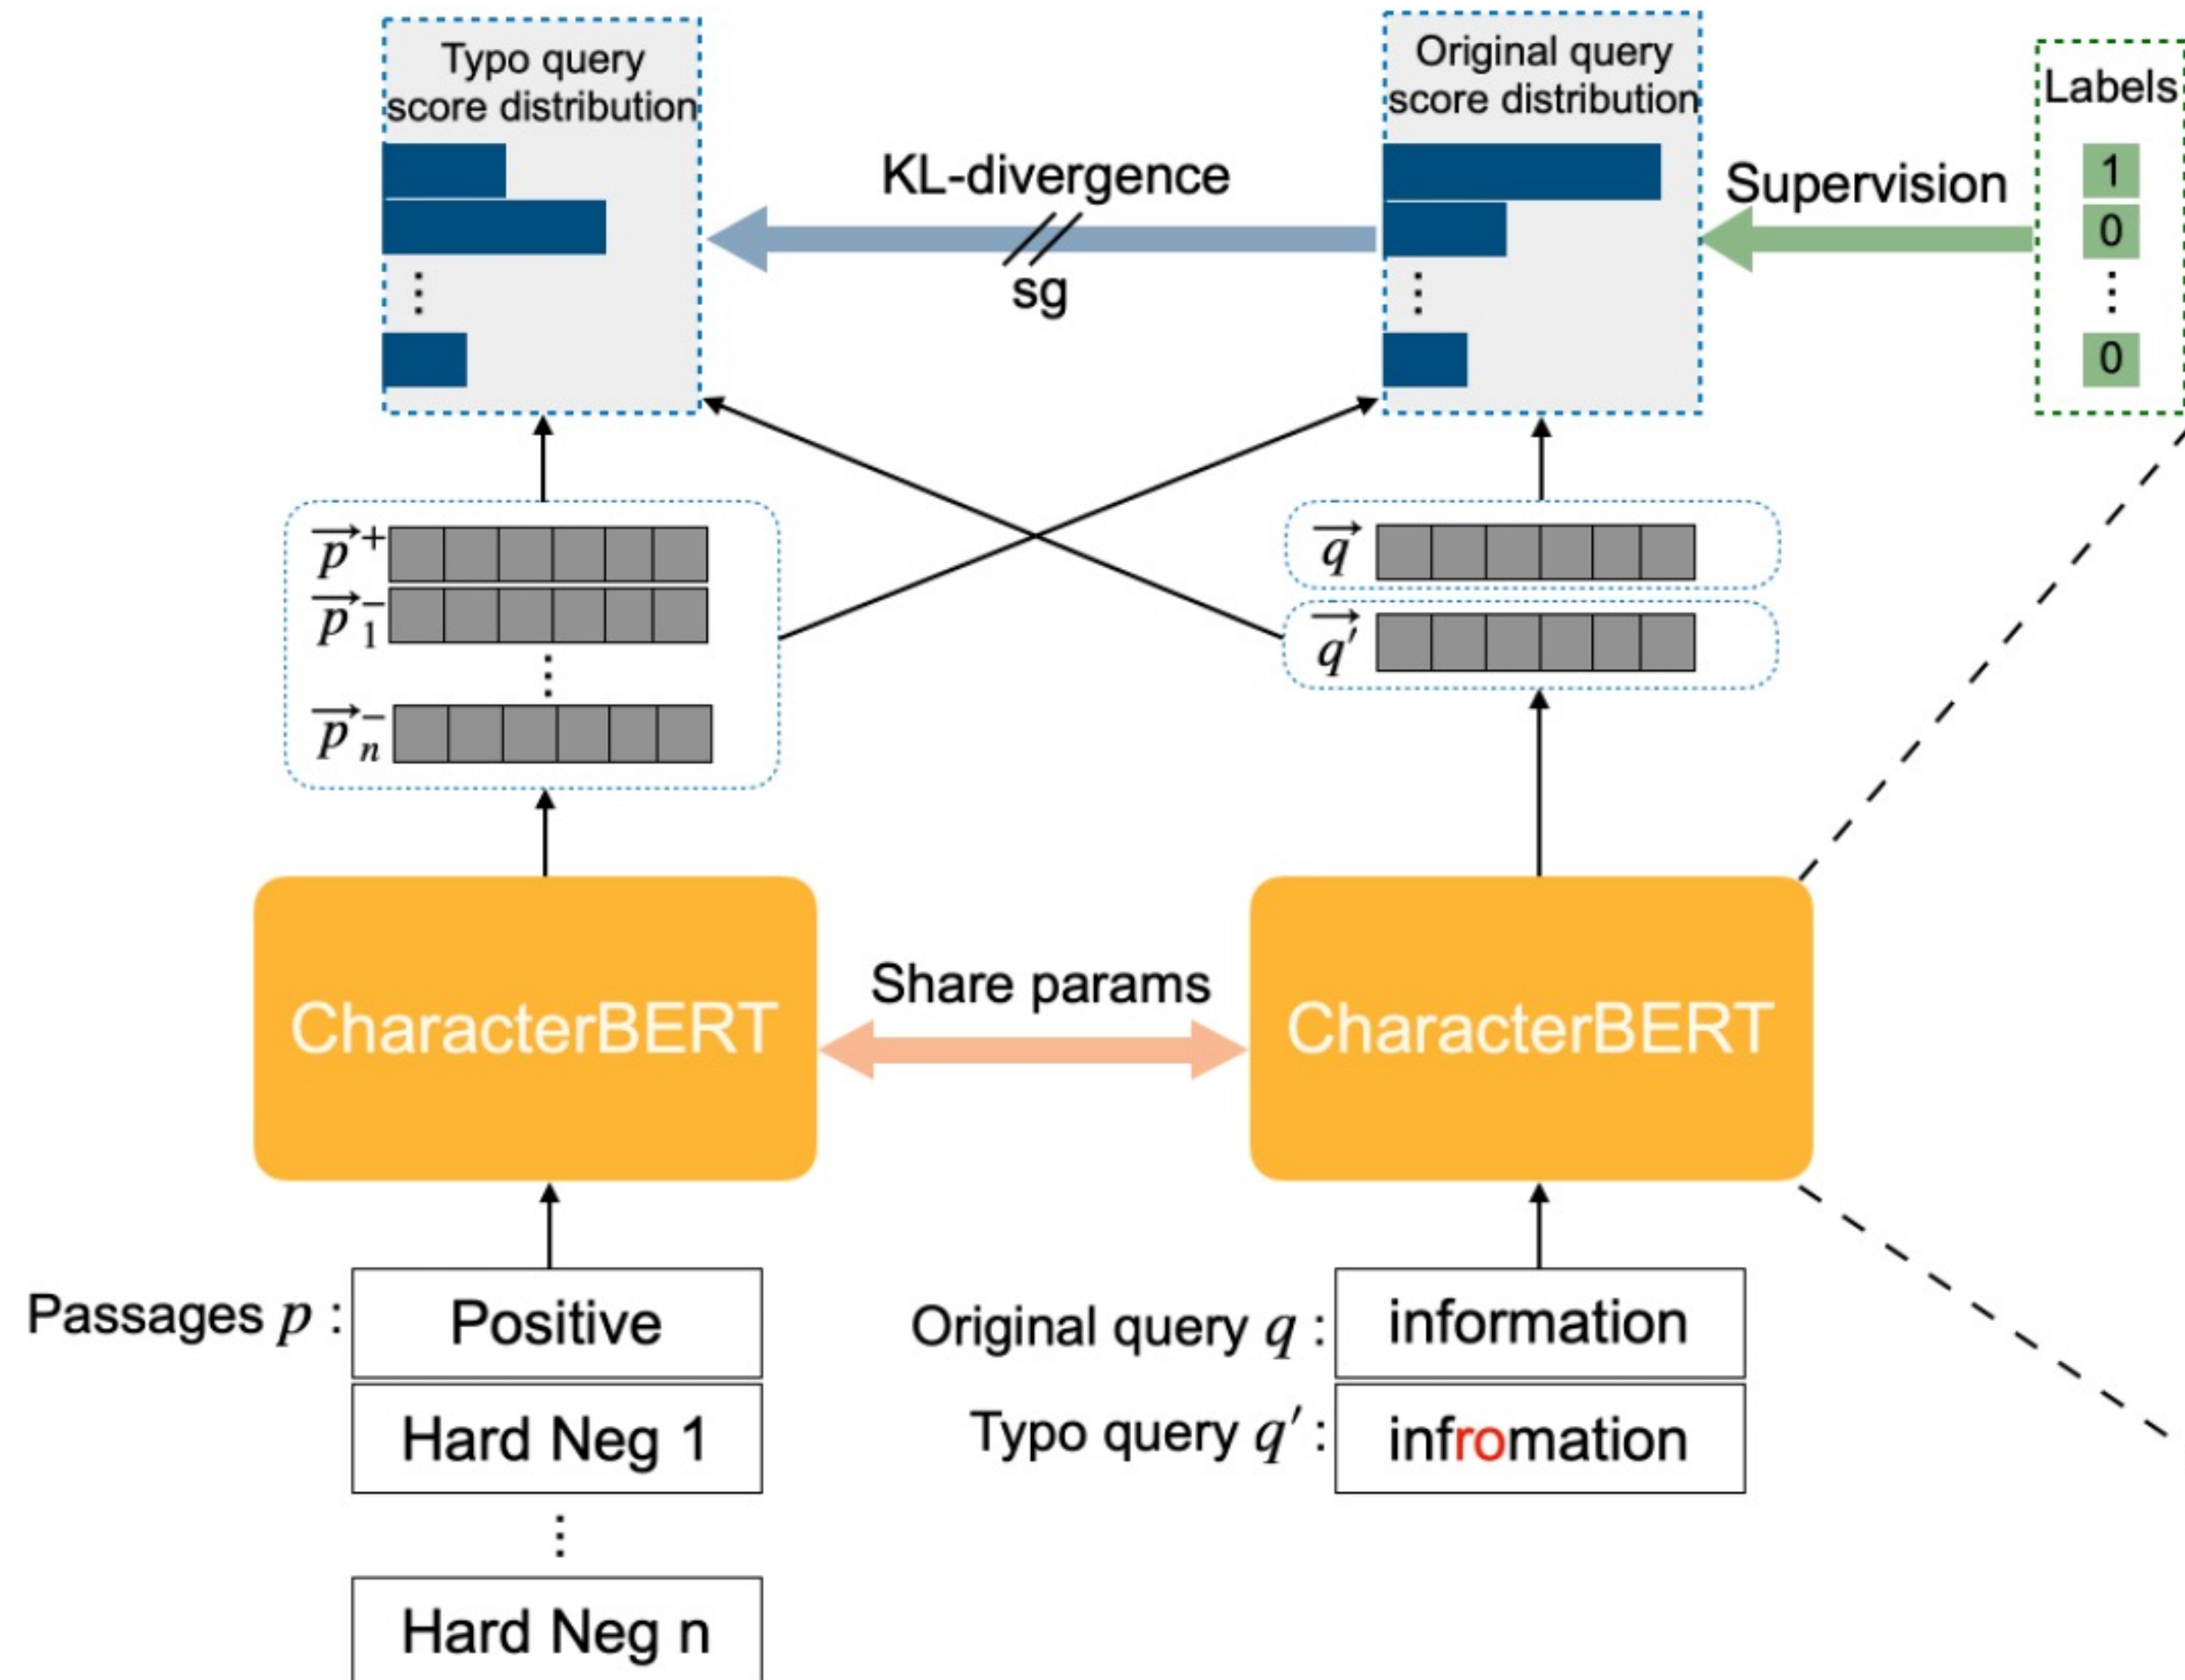
CharacterBERT-DR + Self-Teaching

- Replace BERT WordPiece Tokenizer with CharacterBERT [2] to create query and passages embeddings.
- Does not rely on WordPiece vocabulary: any word will be represented by a single word embedding.

[2] CharacterBERT: Reconciling ELMo and BERT for Word-Level Open-Vocabulary Representations From Characters, Hicham et al, COLING 2020



CharacterBERT-DR + Self-Teaching



1. Make a typo augmentation during training.

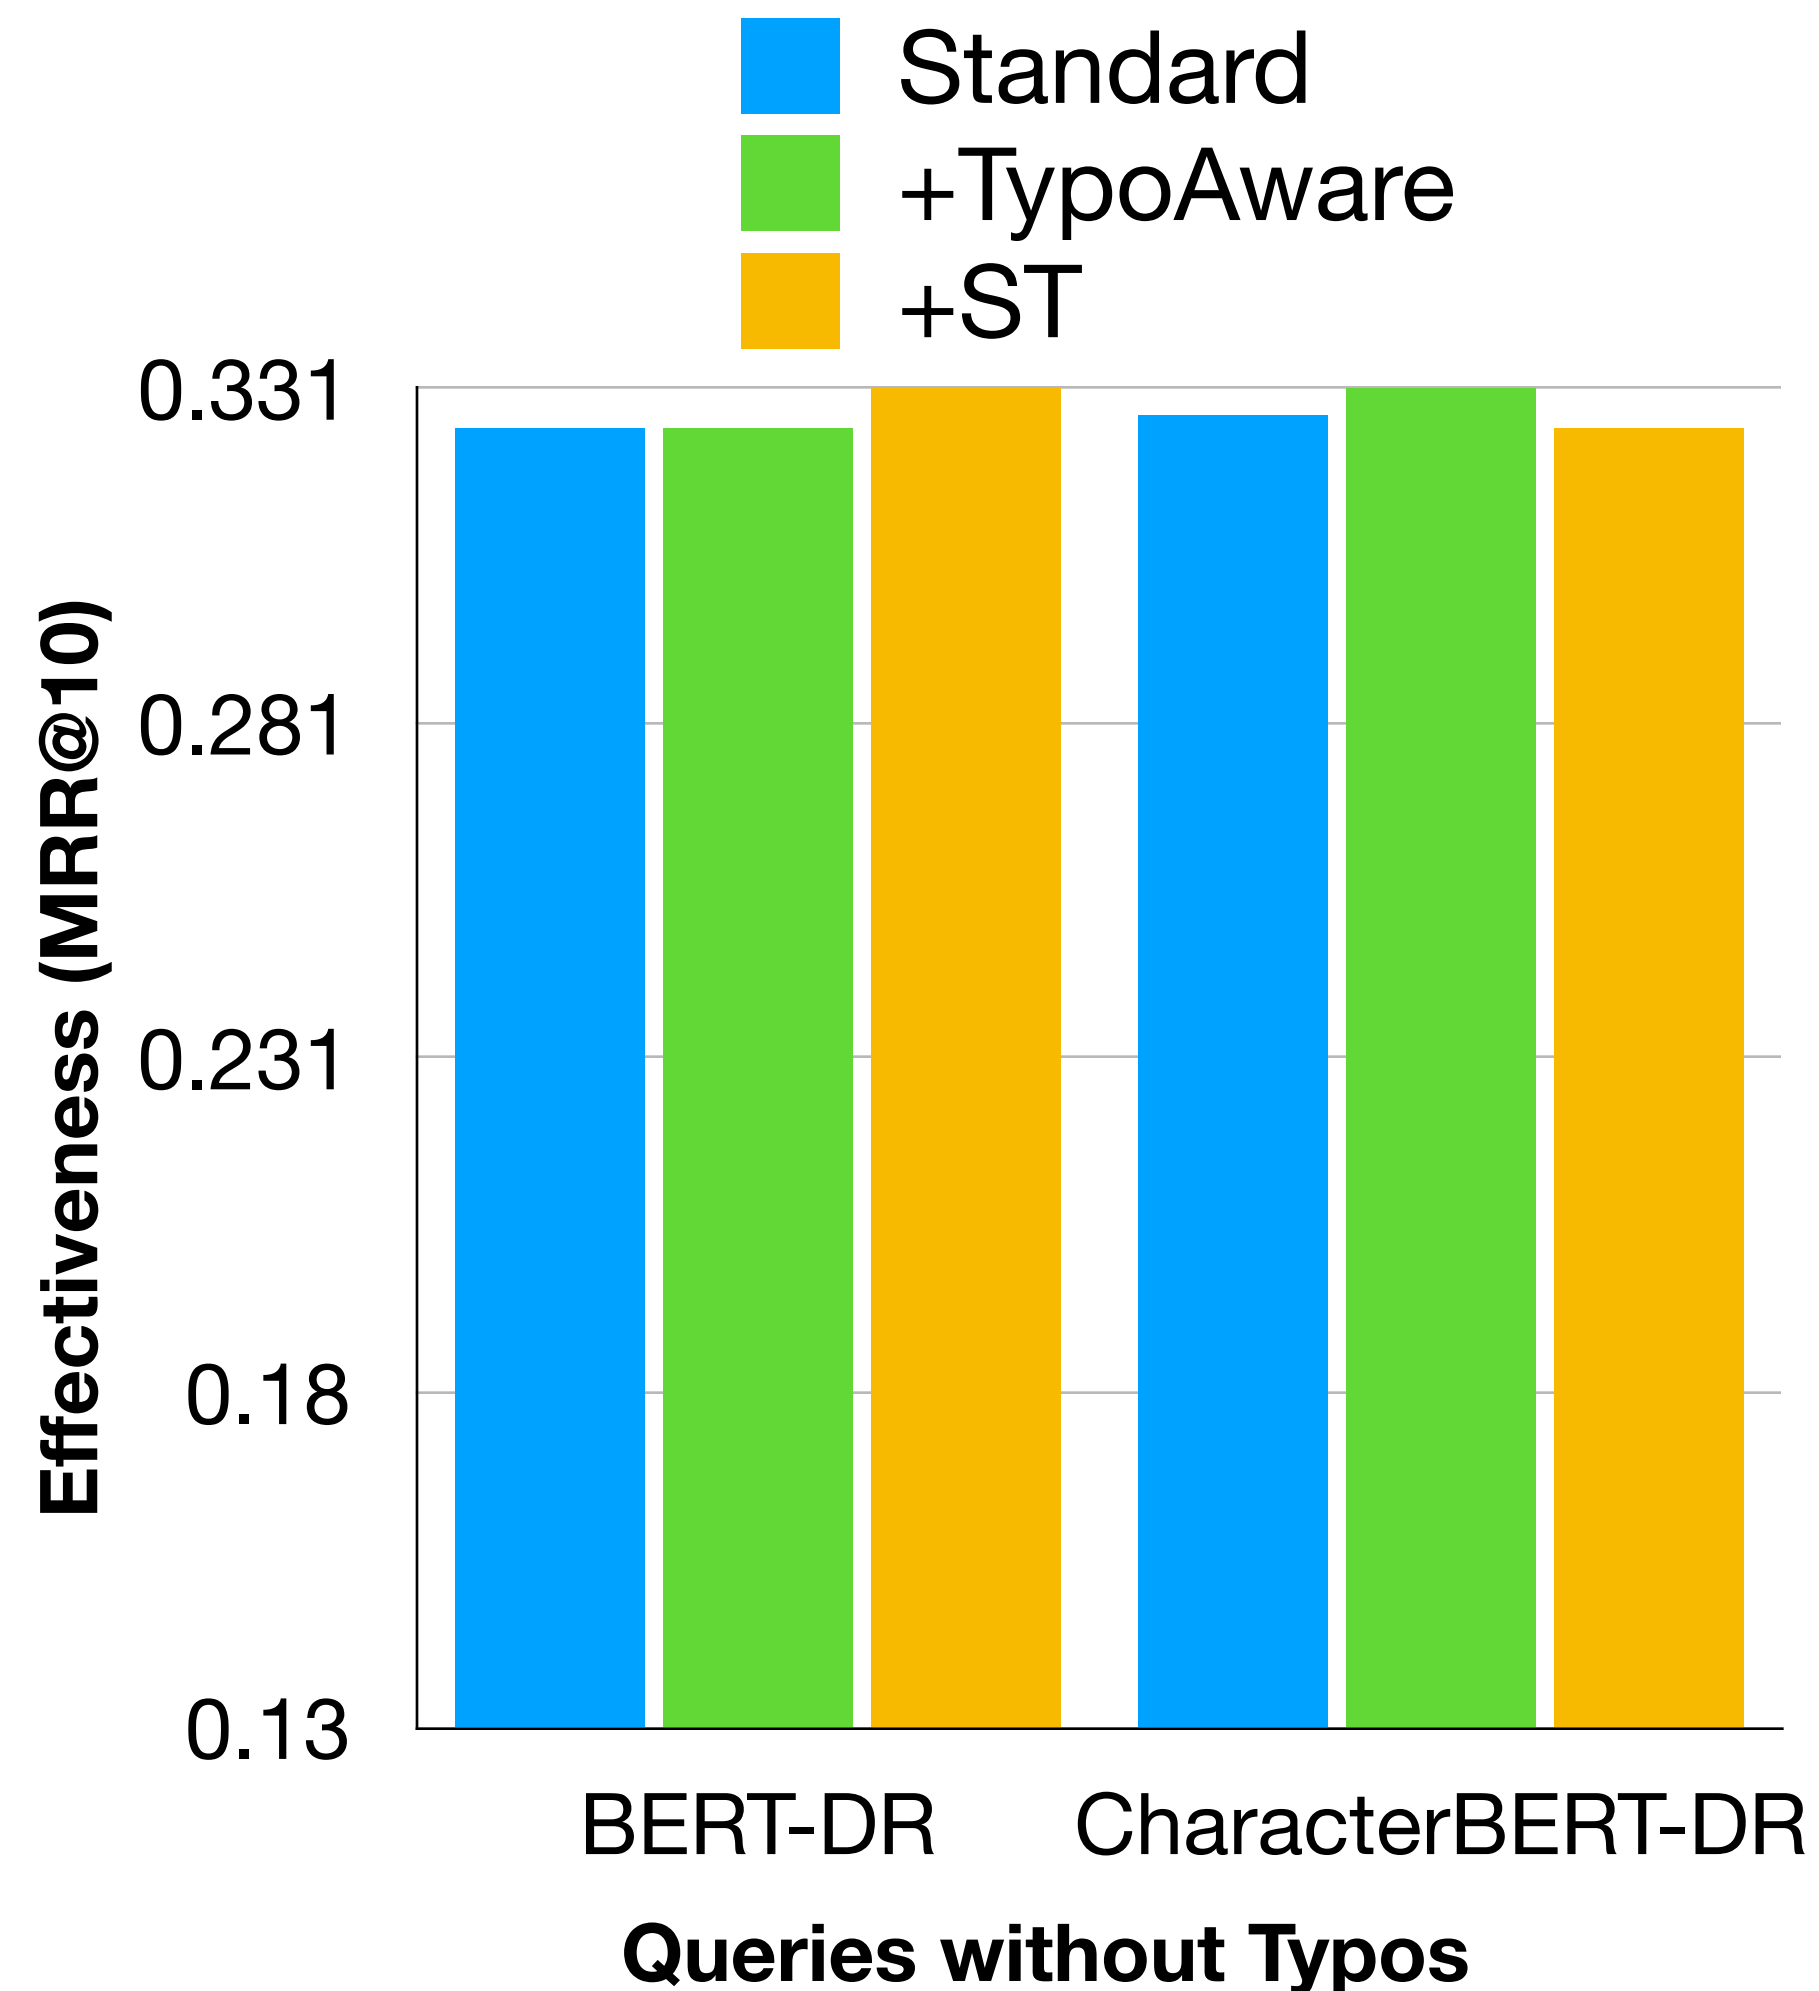
2. Self-Teaching: minimize score distribution difference b/w original query & query with typos:

$$\mathcal{L}_{KL}(\tilde{s}_{q'}, \tilde{s}_q) = \tilde{s}_{q'}(q', p) \cdot \log \frac{\tilde{s}_{q'}(q', p)}{\tilde{s}_q(q, p)}$$

3. Supervised contrastive loss:

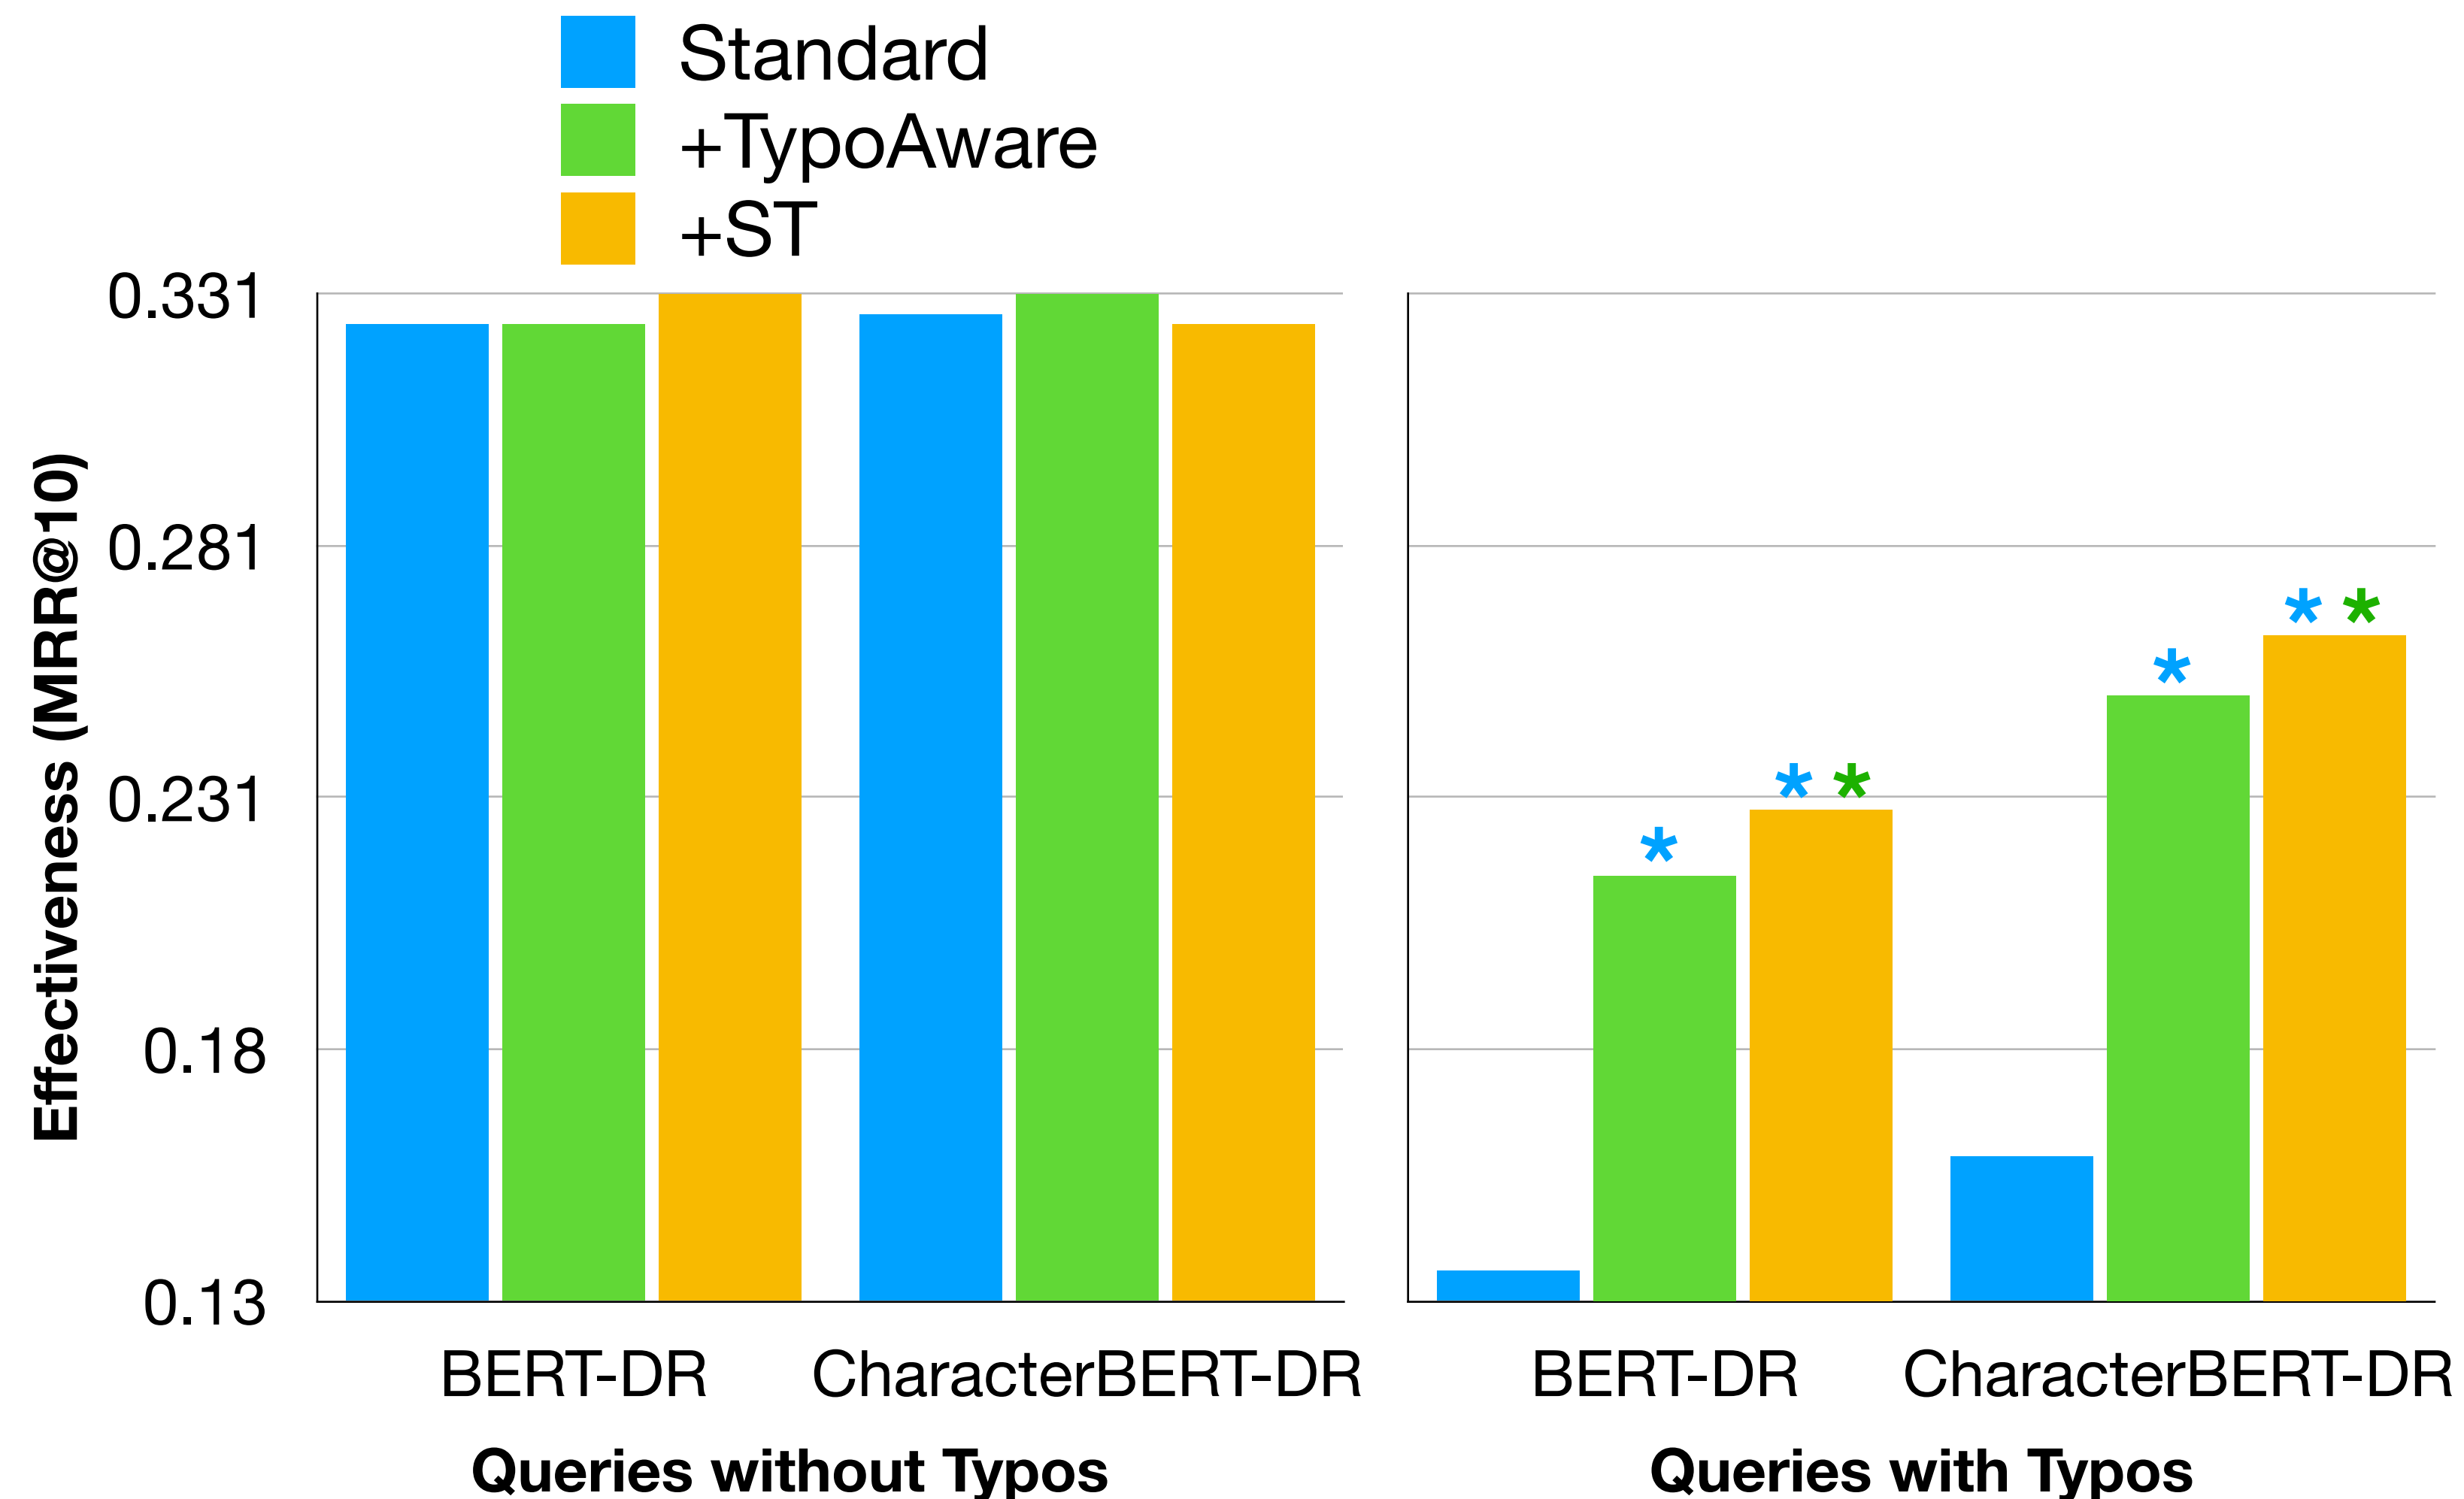
$$\mathcal{L}_{CE}(s_q) = -\log \frac{e^{s_q(q, p^+)}}{e^{s_q(q, p^+)} + \sum_{p^-} e^{s_q(q, p^-)}}$$

Does CharacterBERT+ST produce unwanted effect on queries w/o typos?



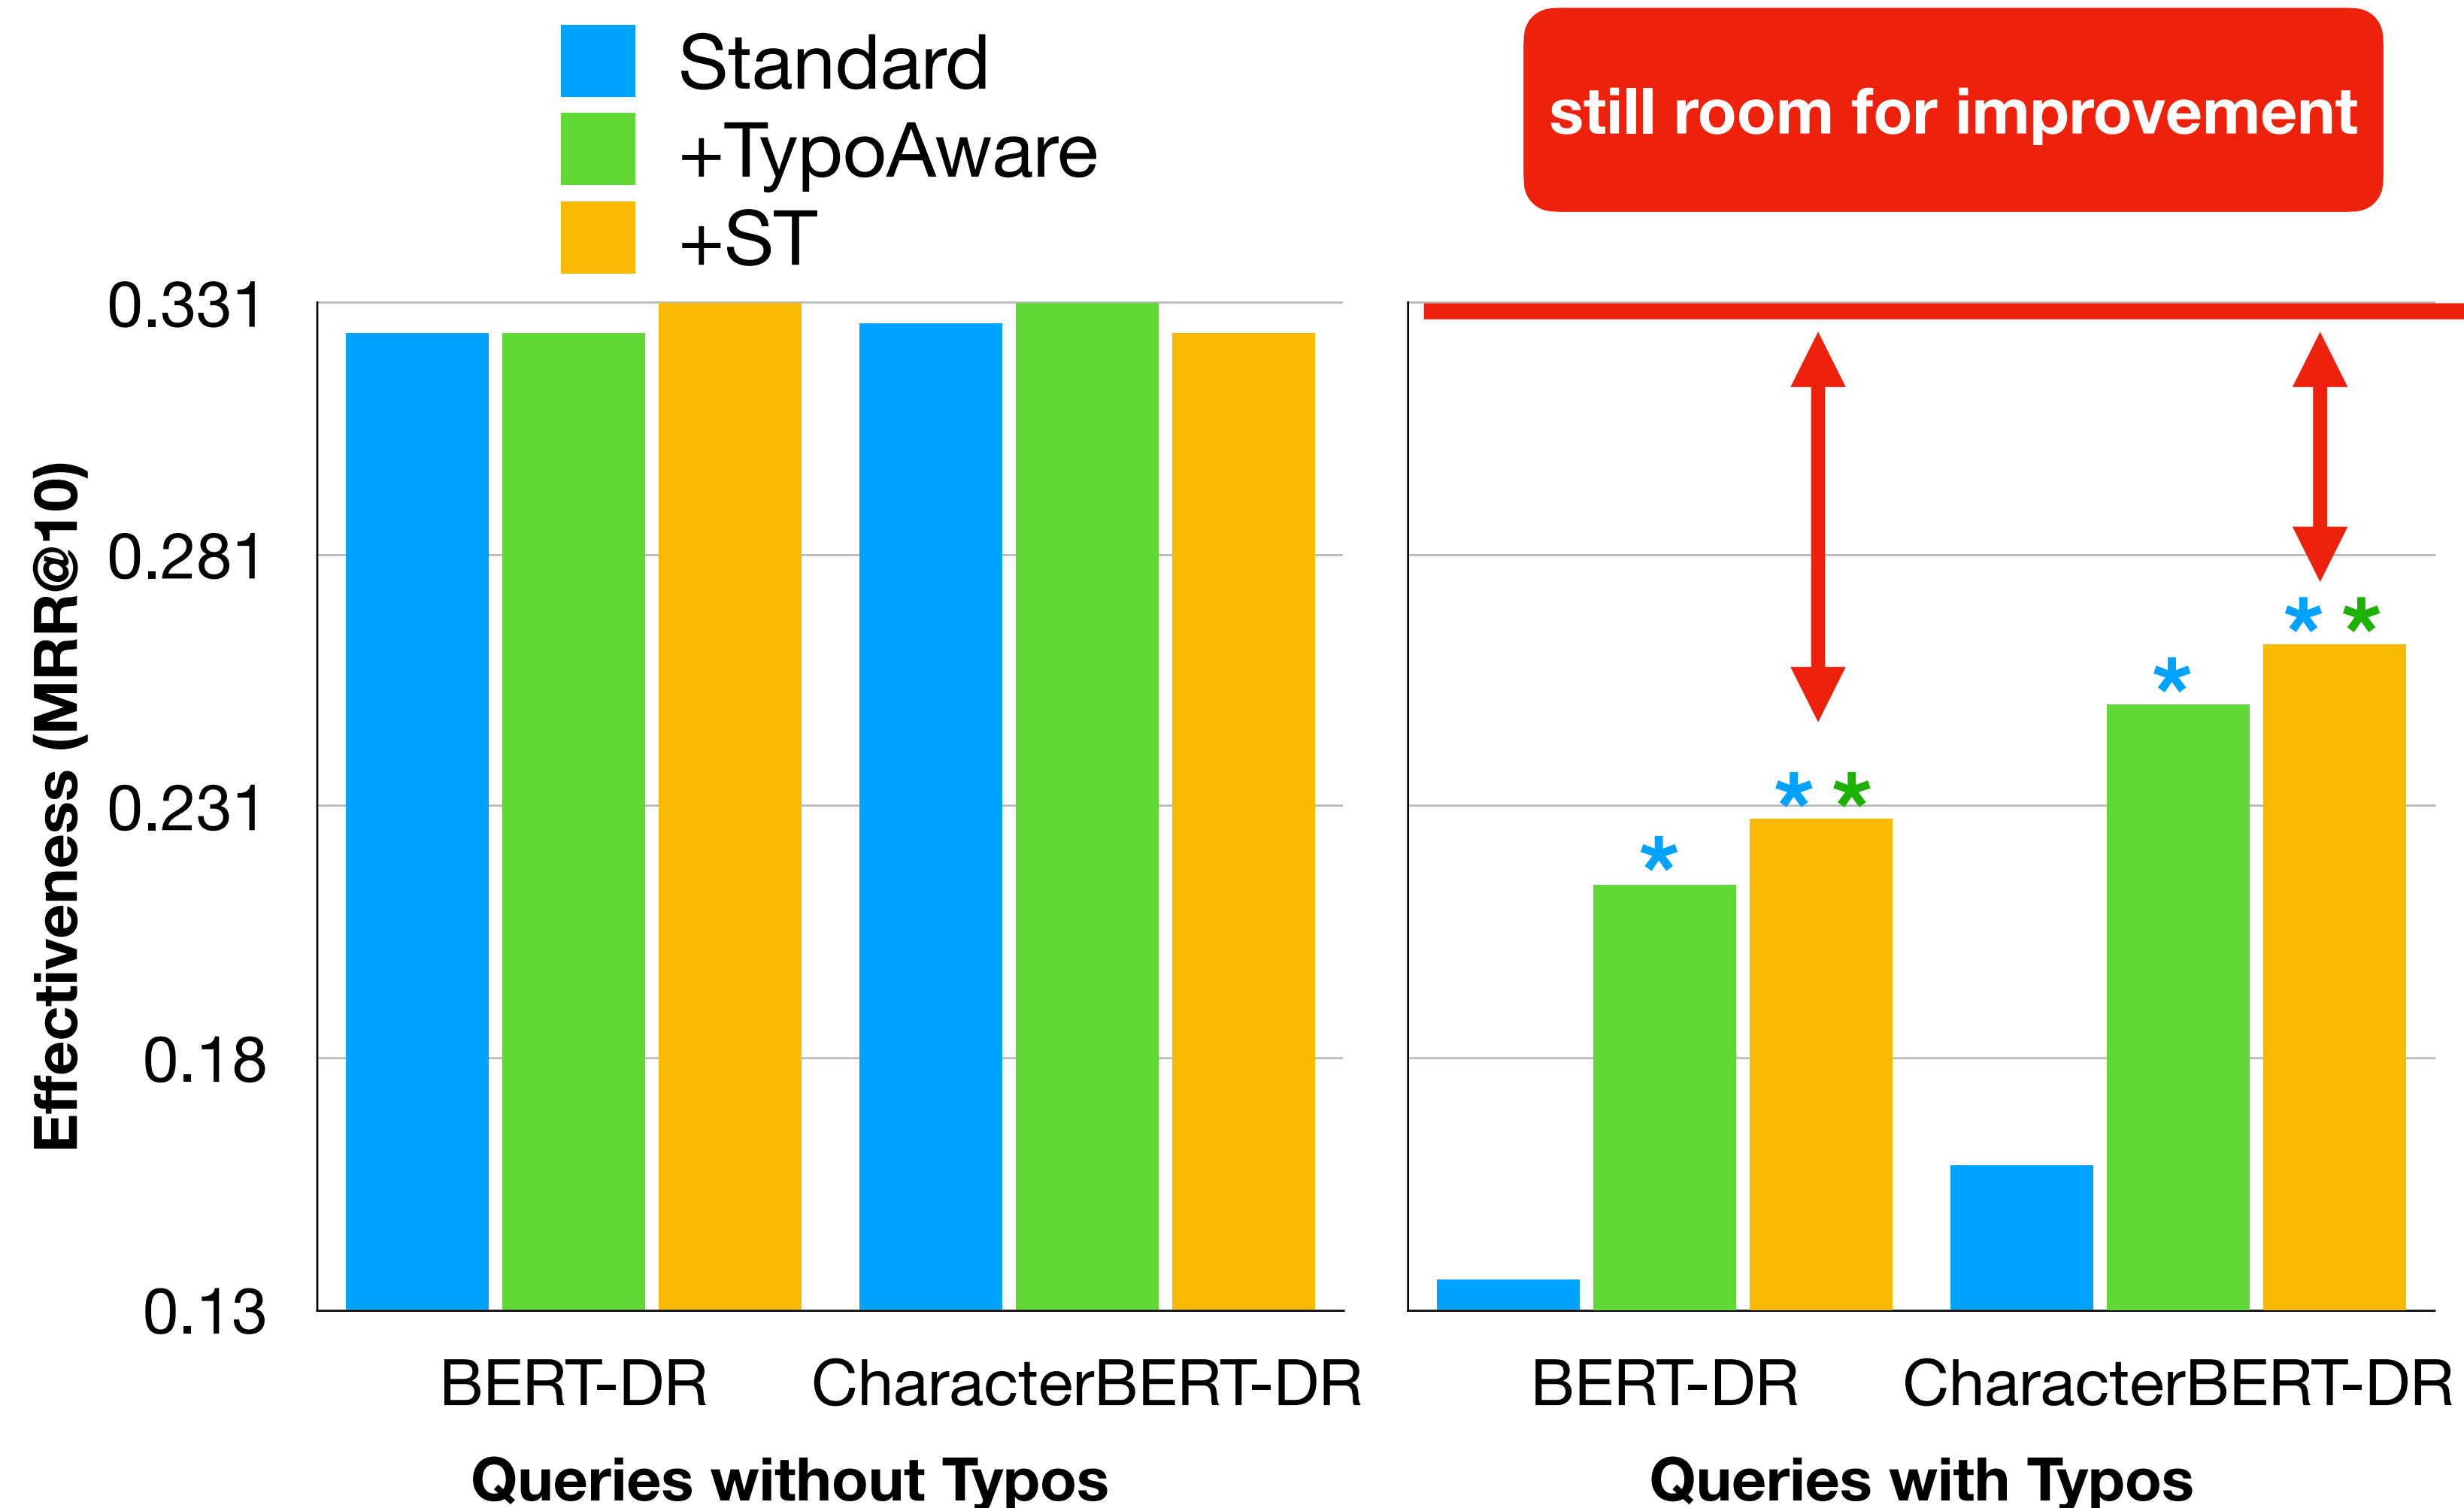
- TypoAware, ST do not provide significant differences on queries without typos: **no risk to use them**

Does CharacterBERT+ST produce improvements on queries with typos?



- TypoAware, ST do not provide significant differences on queries without typos: **no risk to use them**
- **ST provides largest gains** on queries with typos

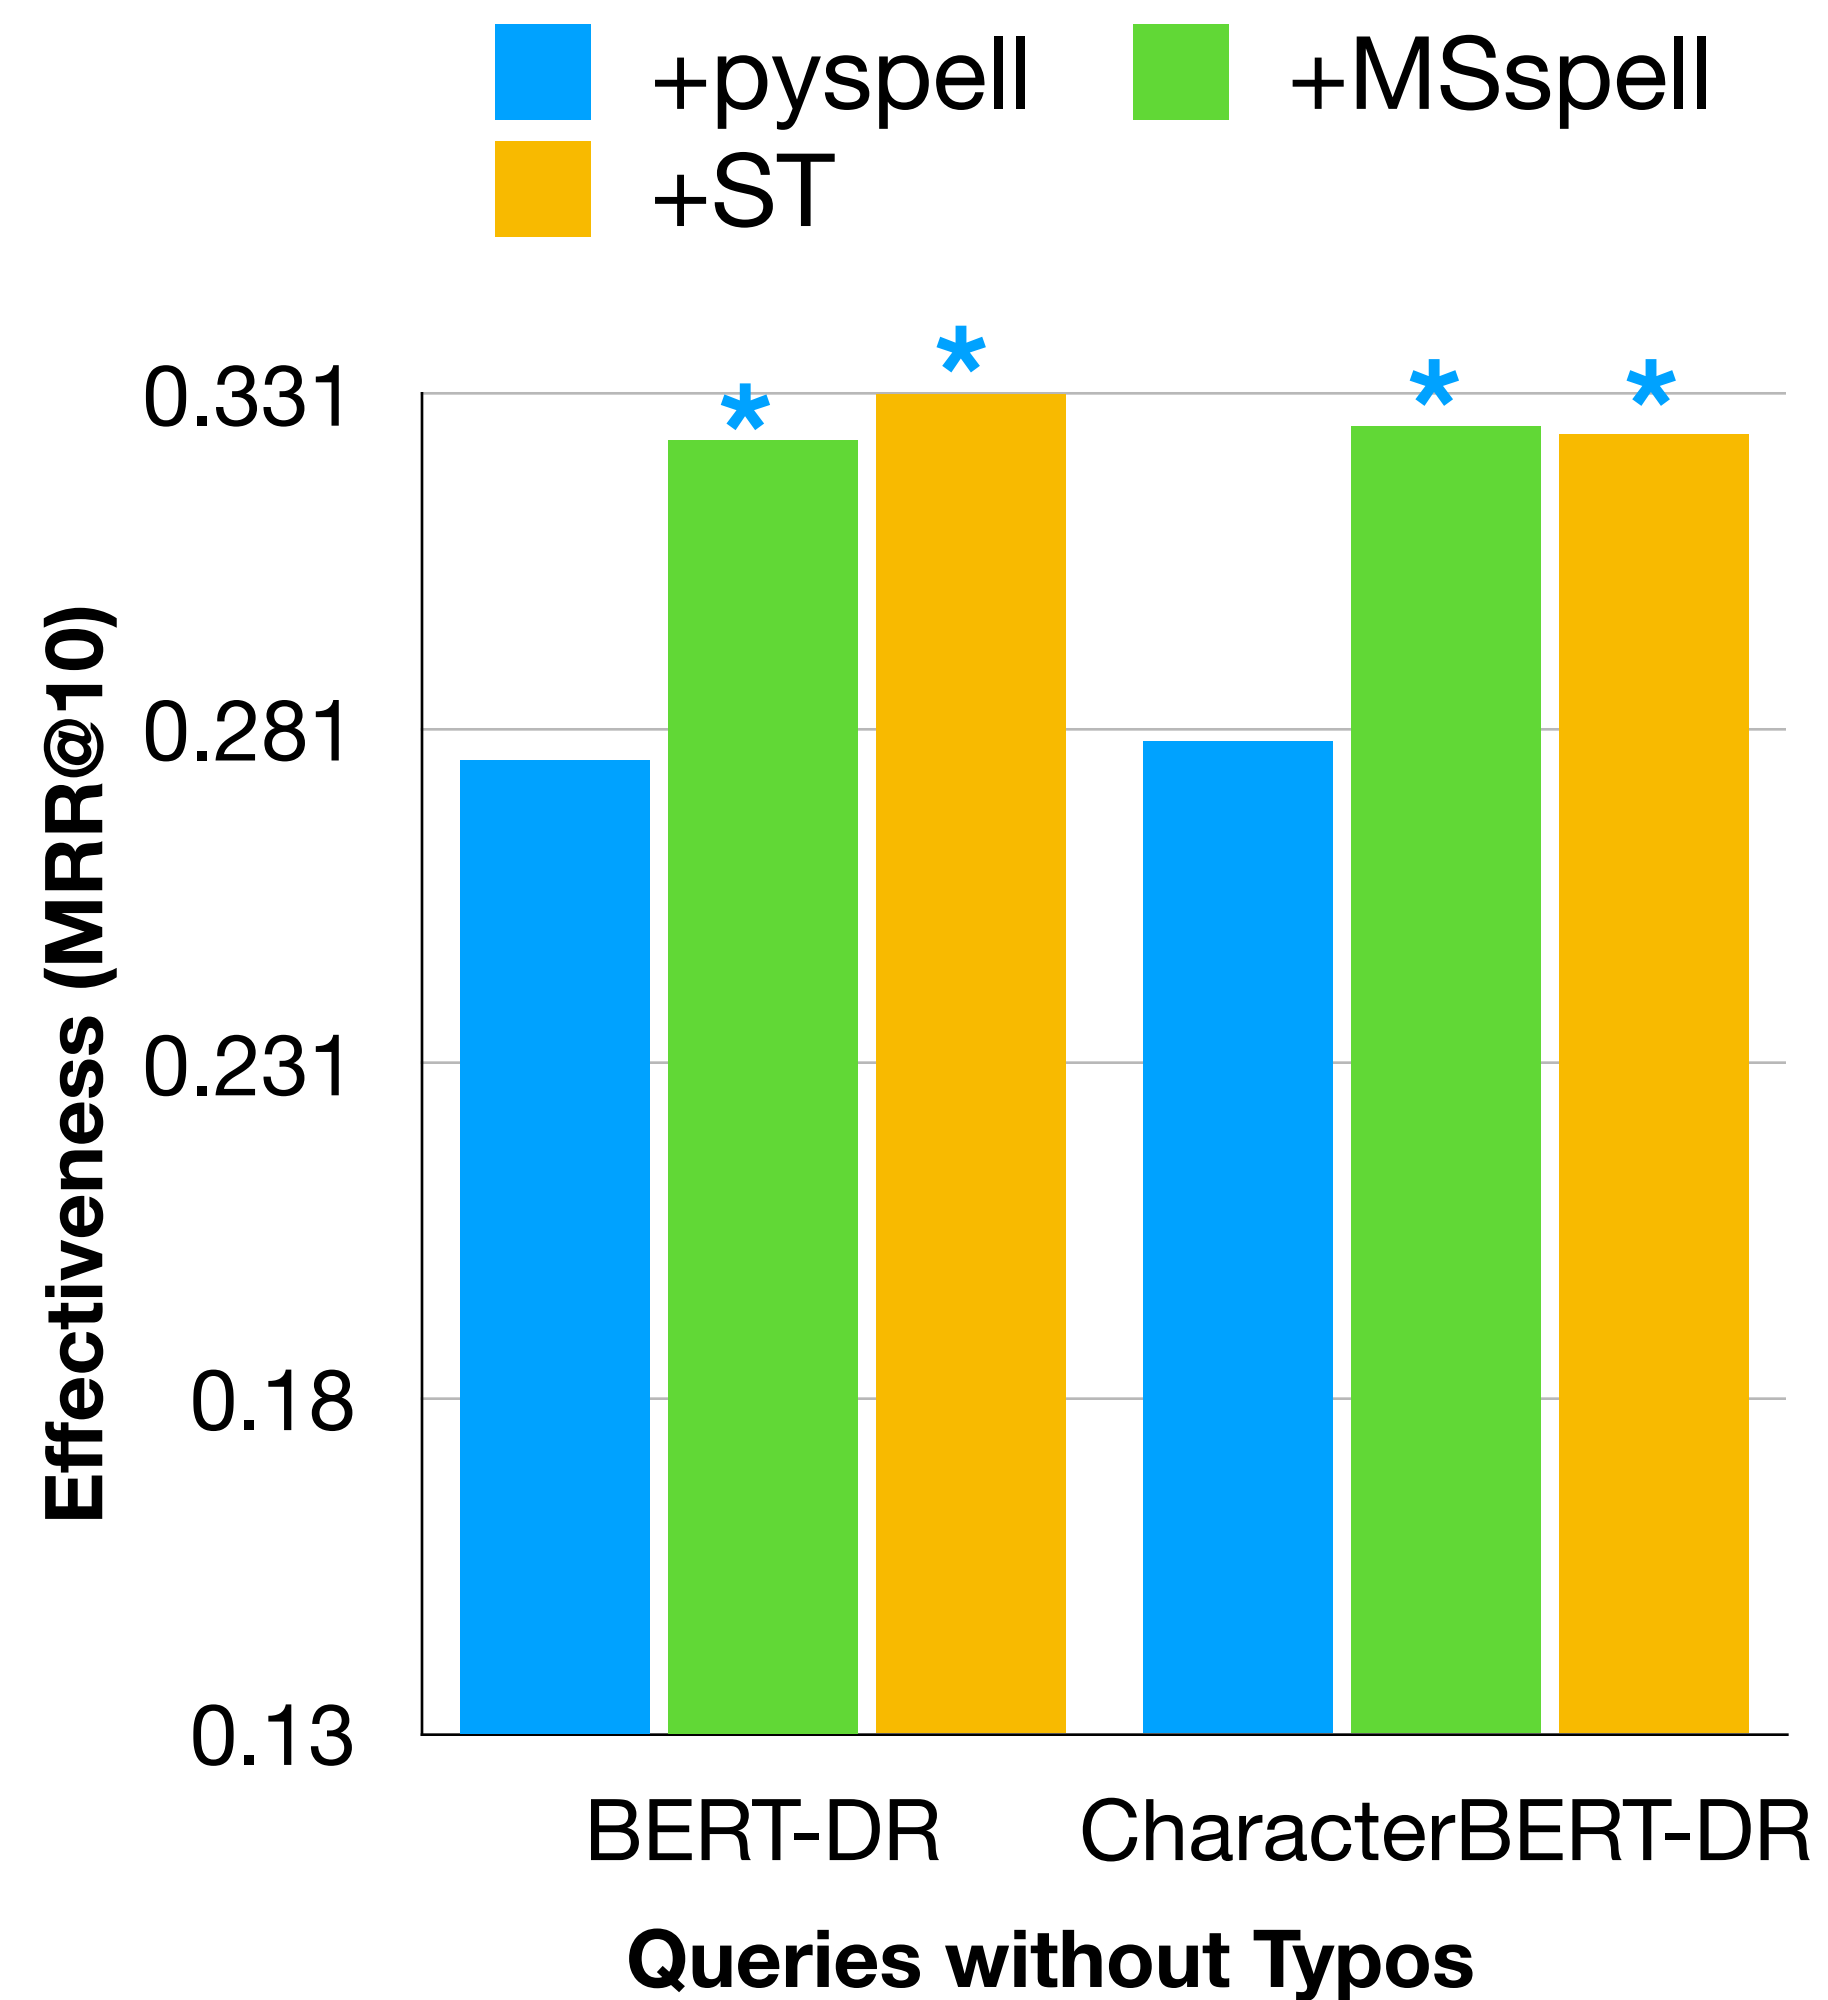
Does CharacterBERT+ST produce improvements on queries with typos?



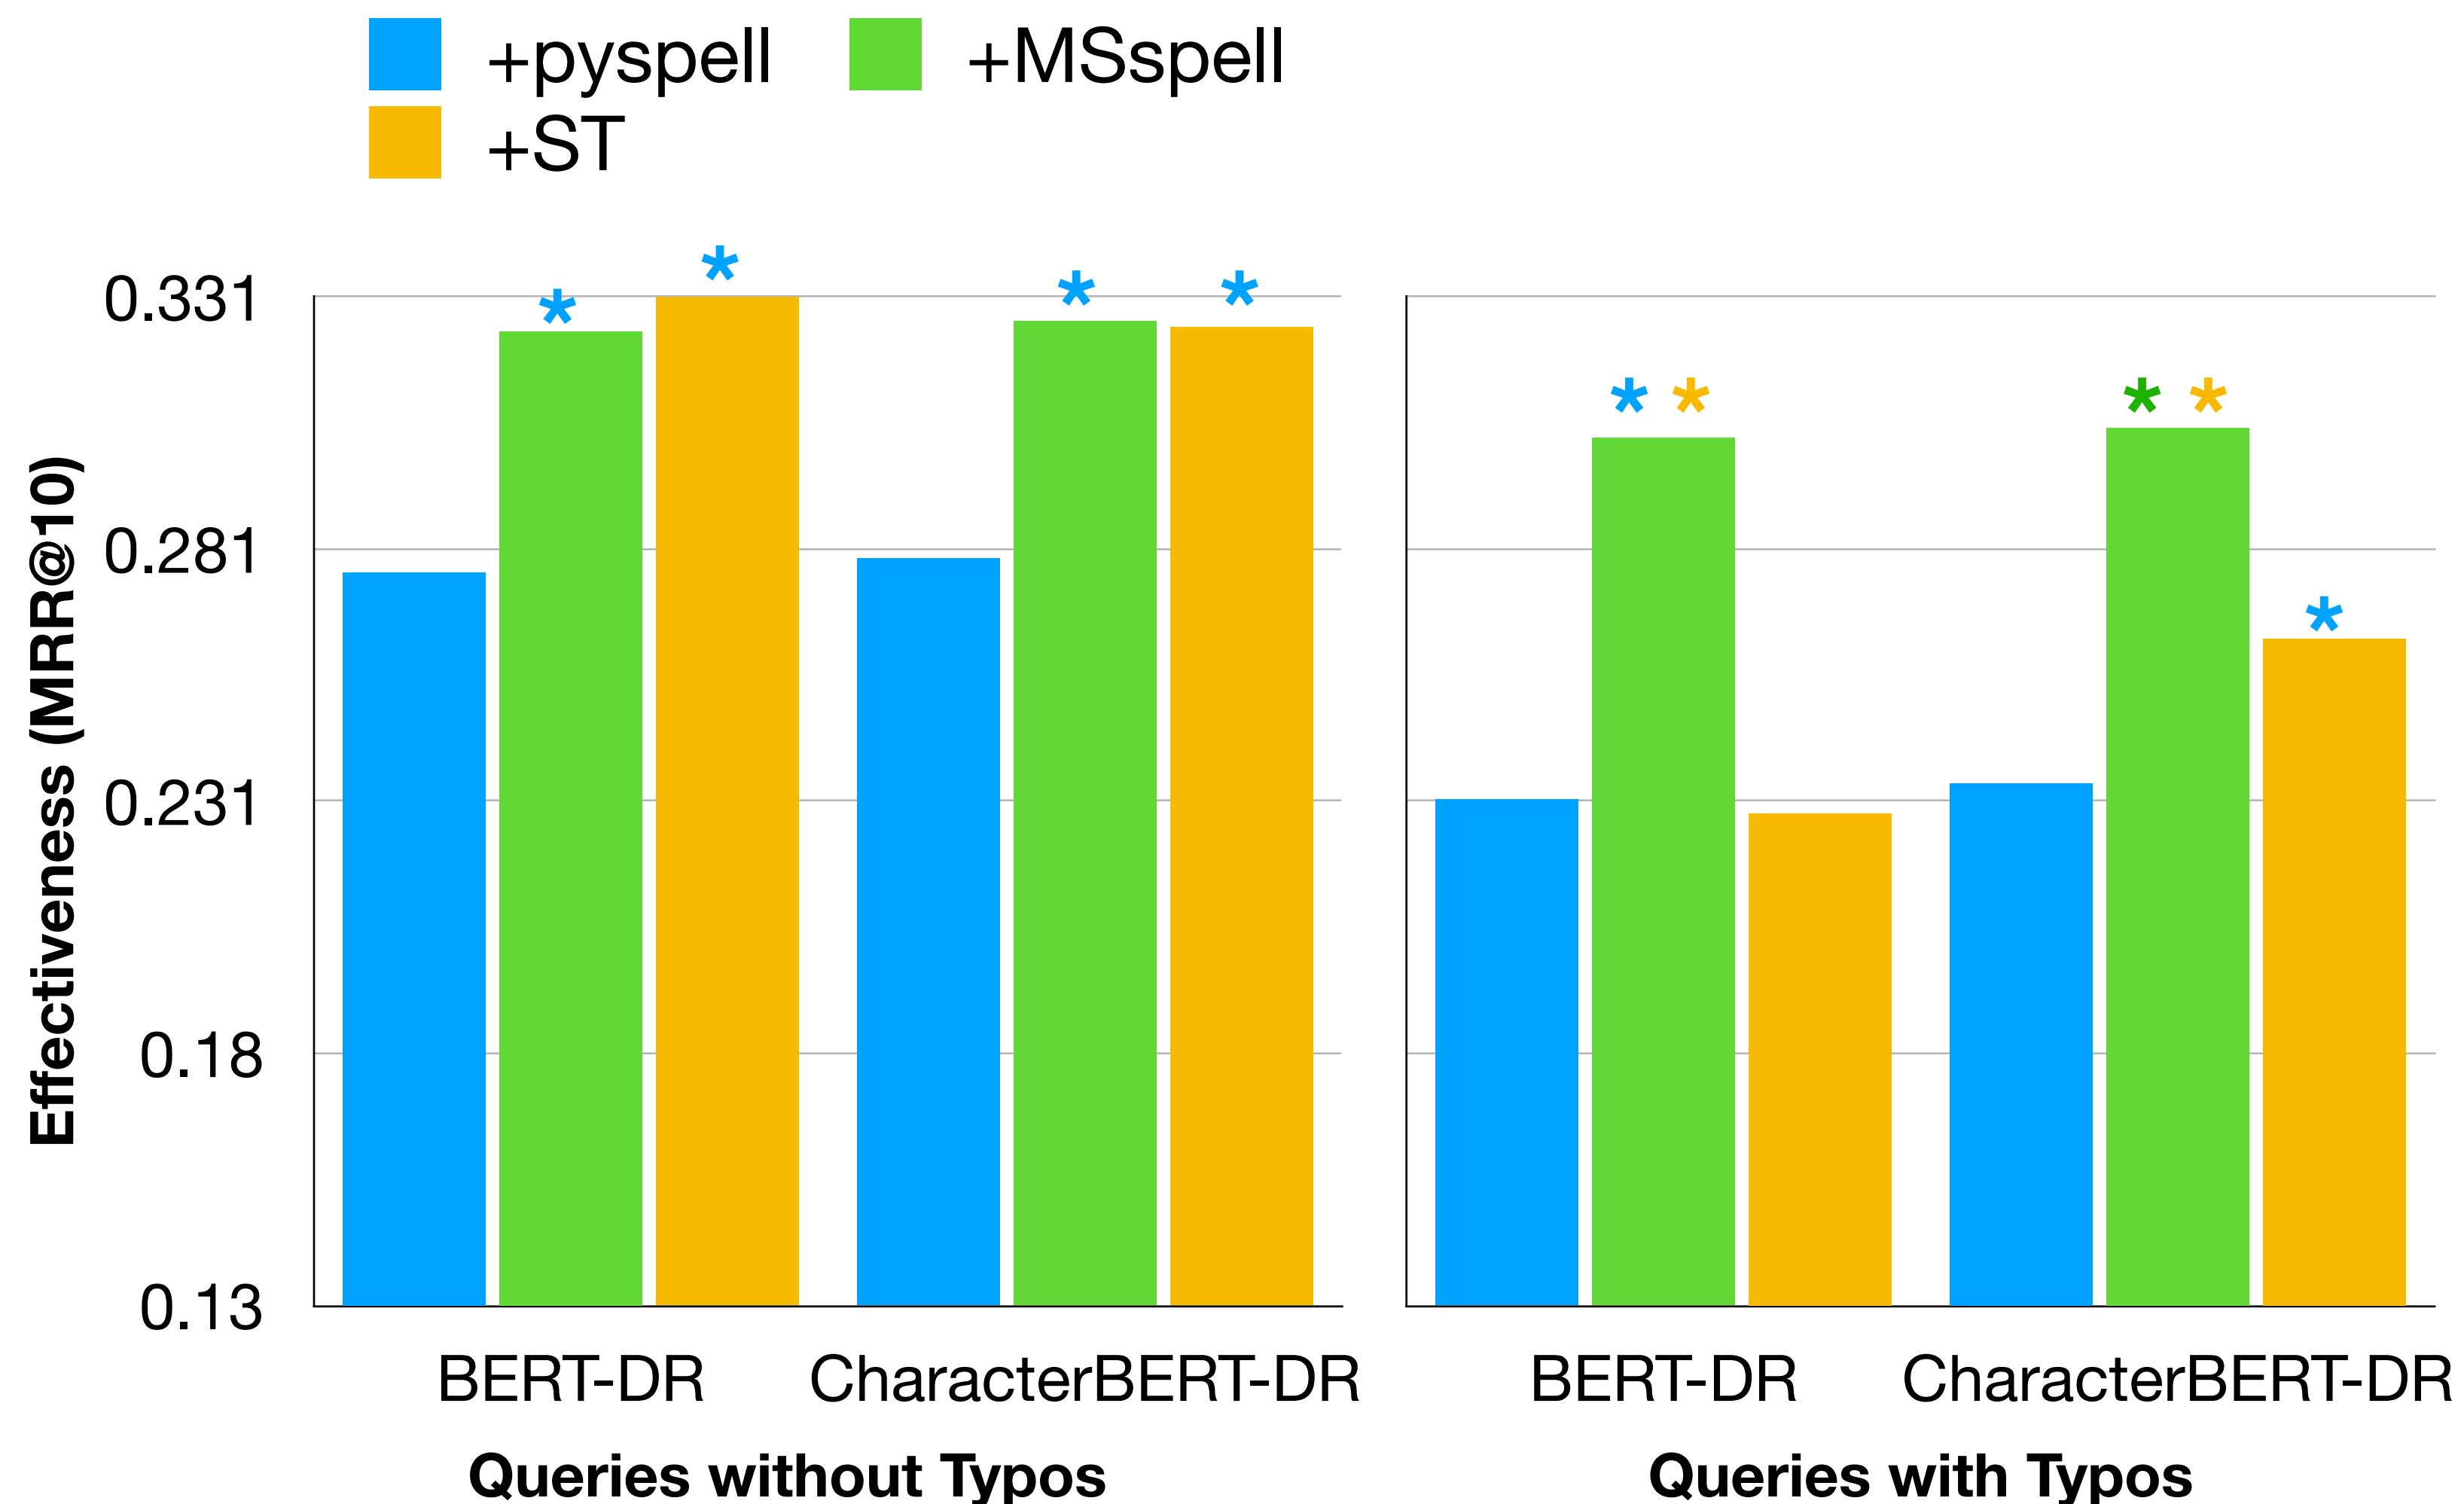
- TypoAware, ST do not provide significant differences on queries without typos: **no risk to use them**
- **ST provides largest gains** on queries with typos

What about using Spell Checkers instead?

What about using Spell Checkers instead?




What about using Spell Checkers instead?



- MSspell provides significantly higher effectiveness
 - Most likely leverages extensive training data
- CharacterBERT+ST better than rule-based spell checker (pyspell)
- Engineering advantages in end-to-end DR pipeline rather than additional spell checker

What about using Spell Checkers instead?



- MSspell provides significantly higher effectiveness
 - Most likely leverages extensive training data
- CharacterBERT+ST better than rule-based spell checker (pyspell)
- Engineering advantages in end-to-end DR pipeline rather than additional spell checker

Take-aways

- Typical Dense Retrievers do not perform well on queries with typos. Augmentation at training (EMNLP 2021) only goes so far
- The **key issue is in how a word with typos vs. without-typo is represented**
 - Bringing these two representations closer improves effectiveness
 - BERT's Tokenizer major source for representation differences
- Replacing tokenizer with CharacterBERT encoder and using ST to further bring representations close drastically improves robustness of Dense Retrievers
- We also provide a **new dataset** for evaluation with real queries with typos



Additional Material

CharacterBERT and Self-Teaching for Improving the Robustness of Dense Retrievers on Queries with Typos

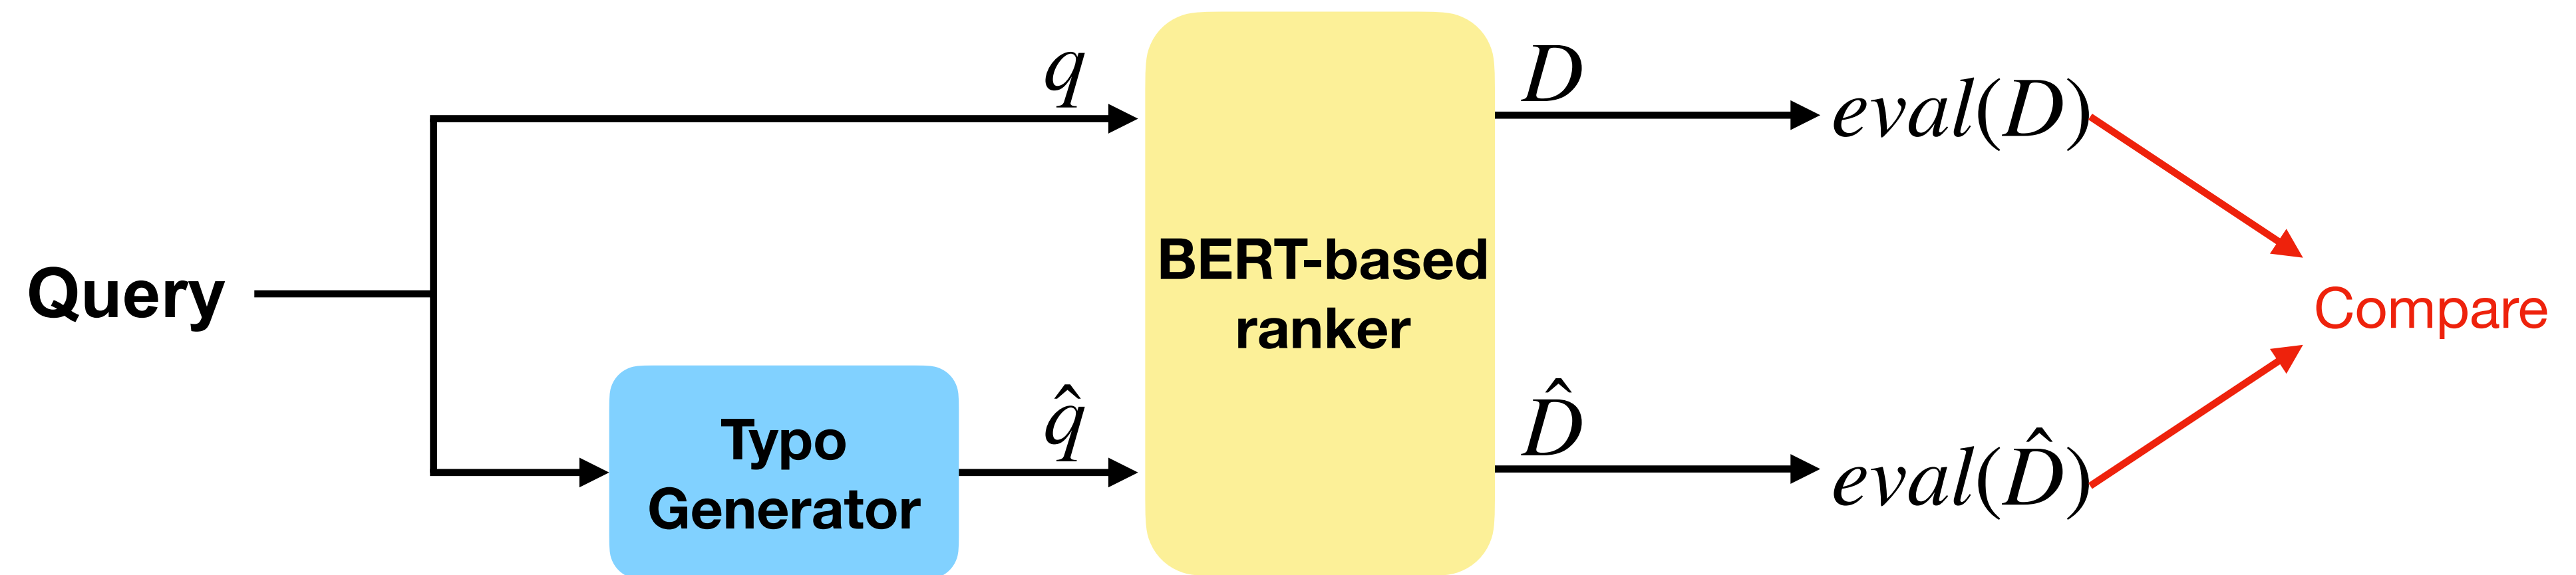
Shengyao Zhuang & Guido Zuccon

{s.zhuang,g.zuccon}@uq.edu.au

ielab, The University of Queensland, Australia

www.ielab.io

Evaluation with typo generator

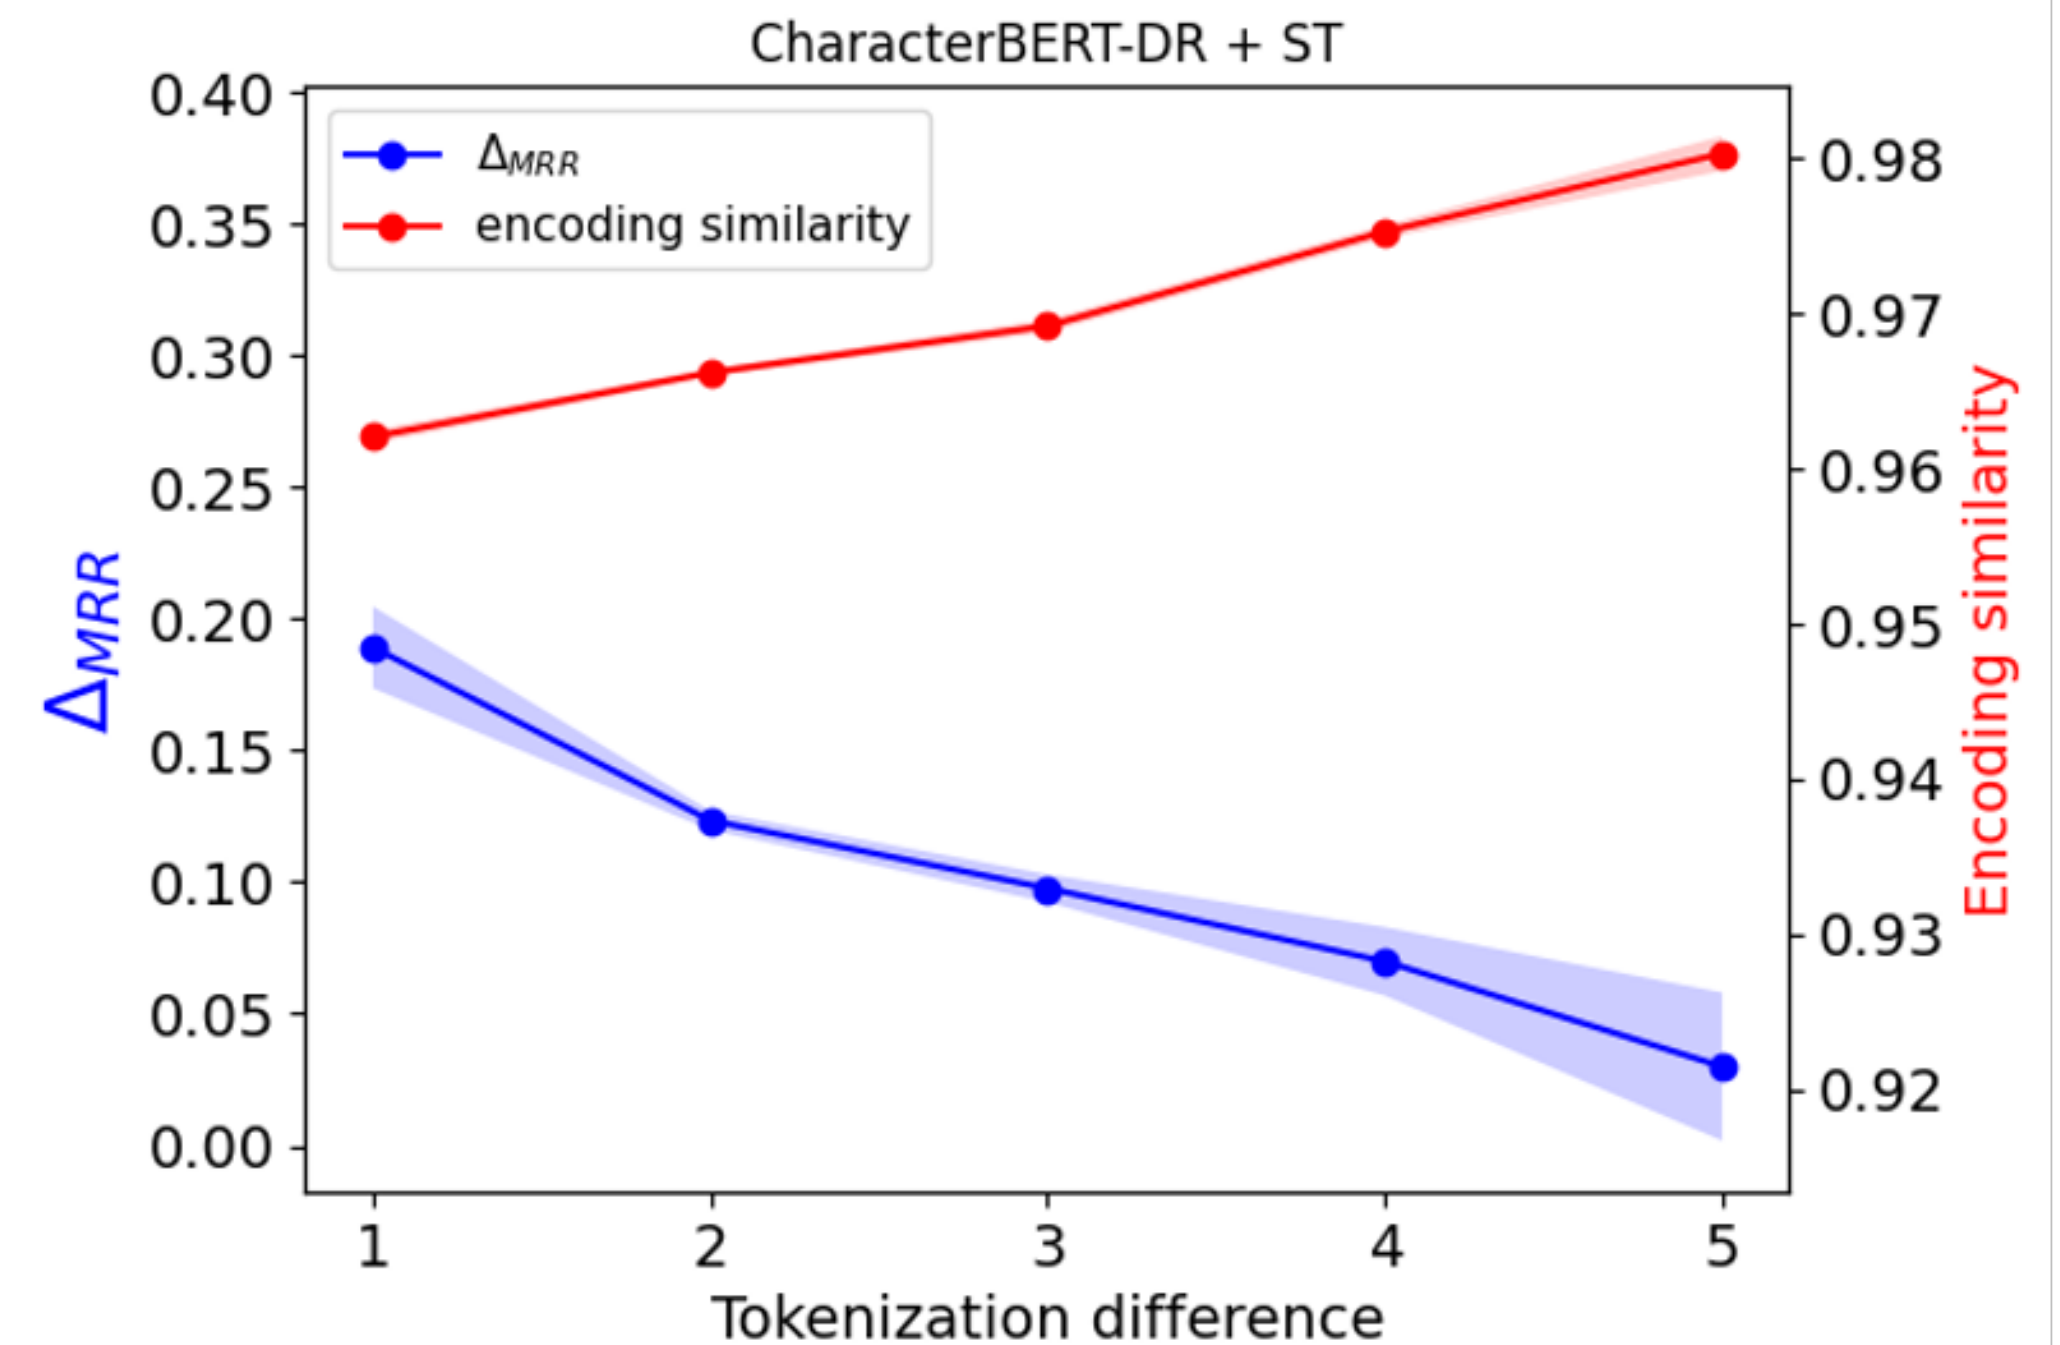
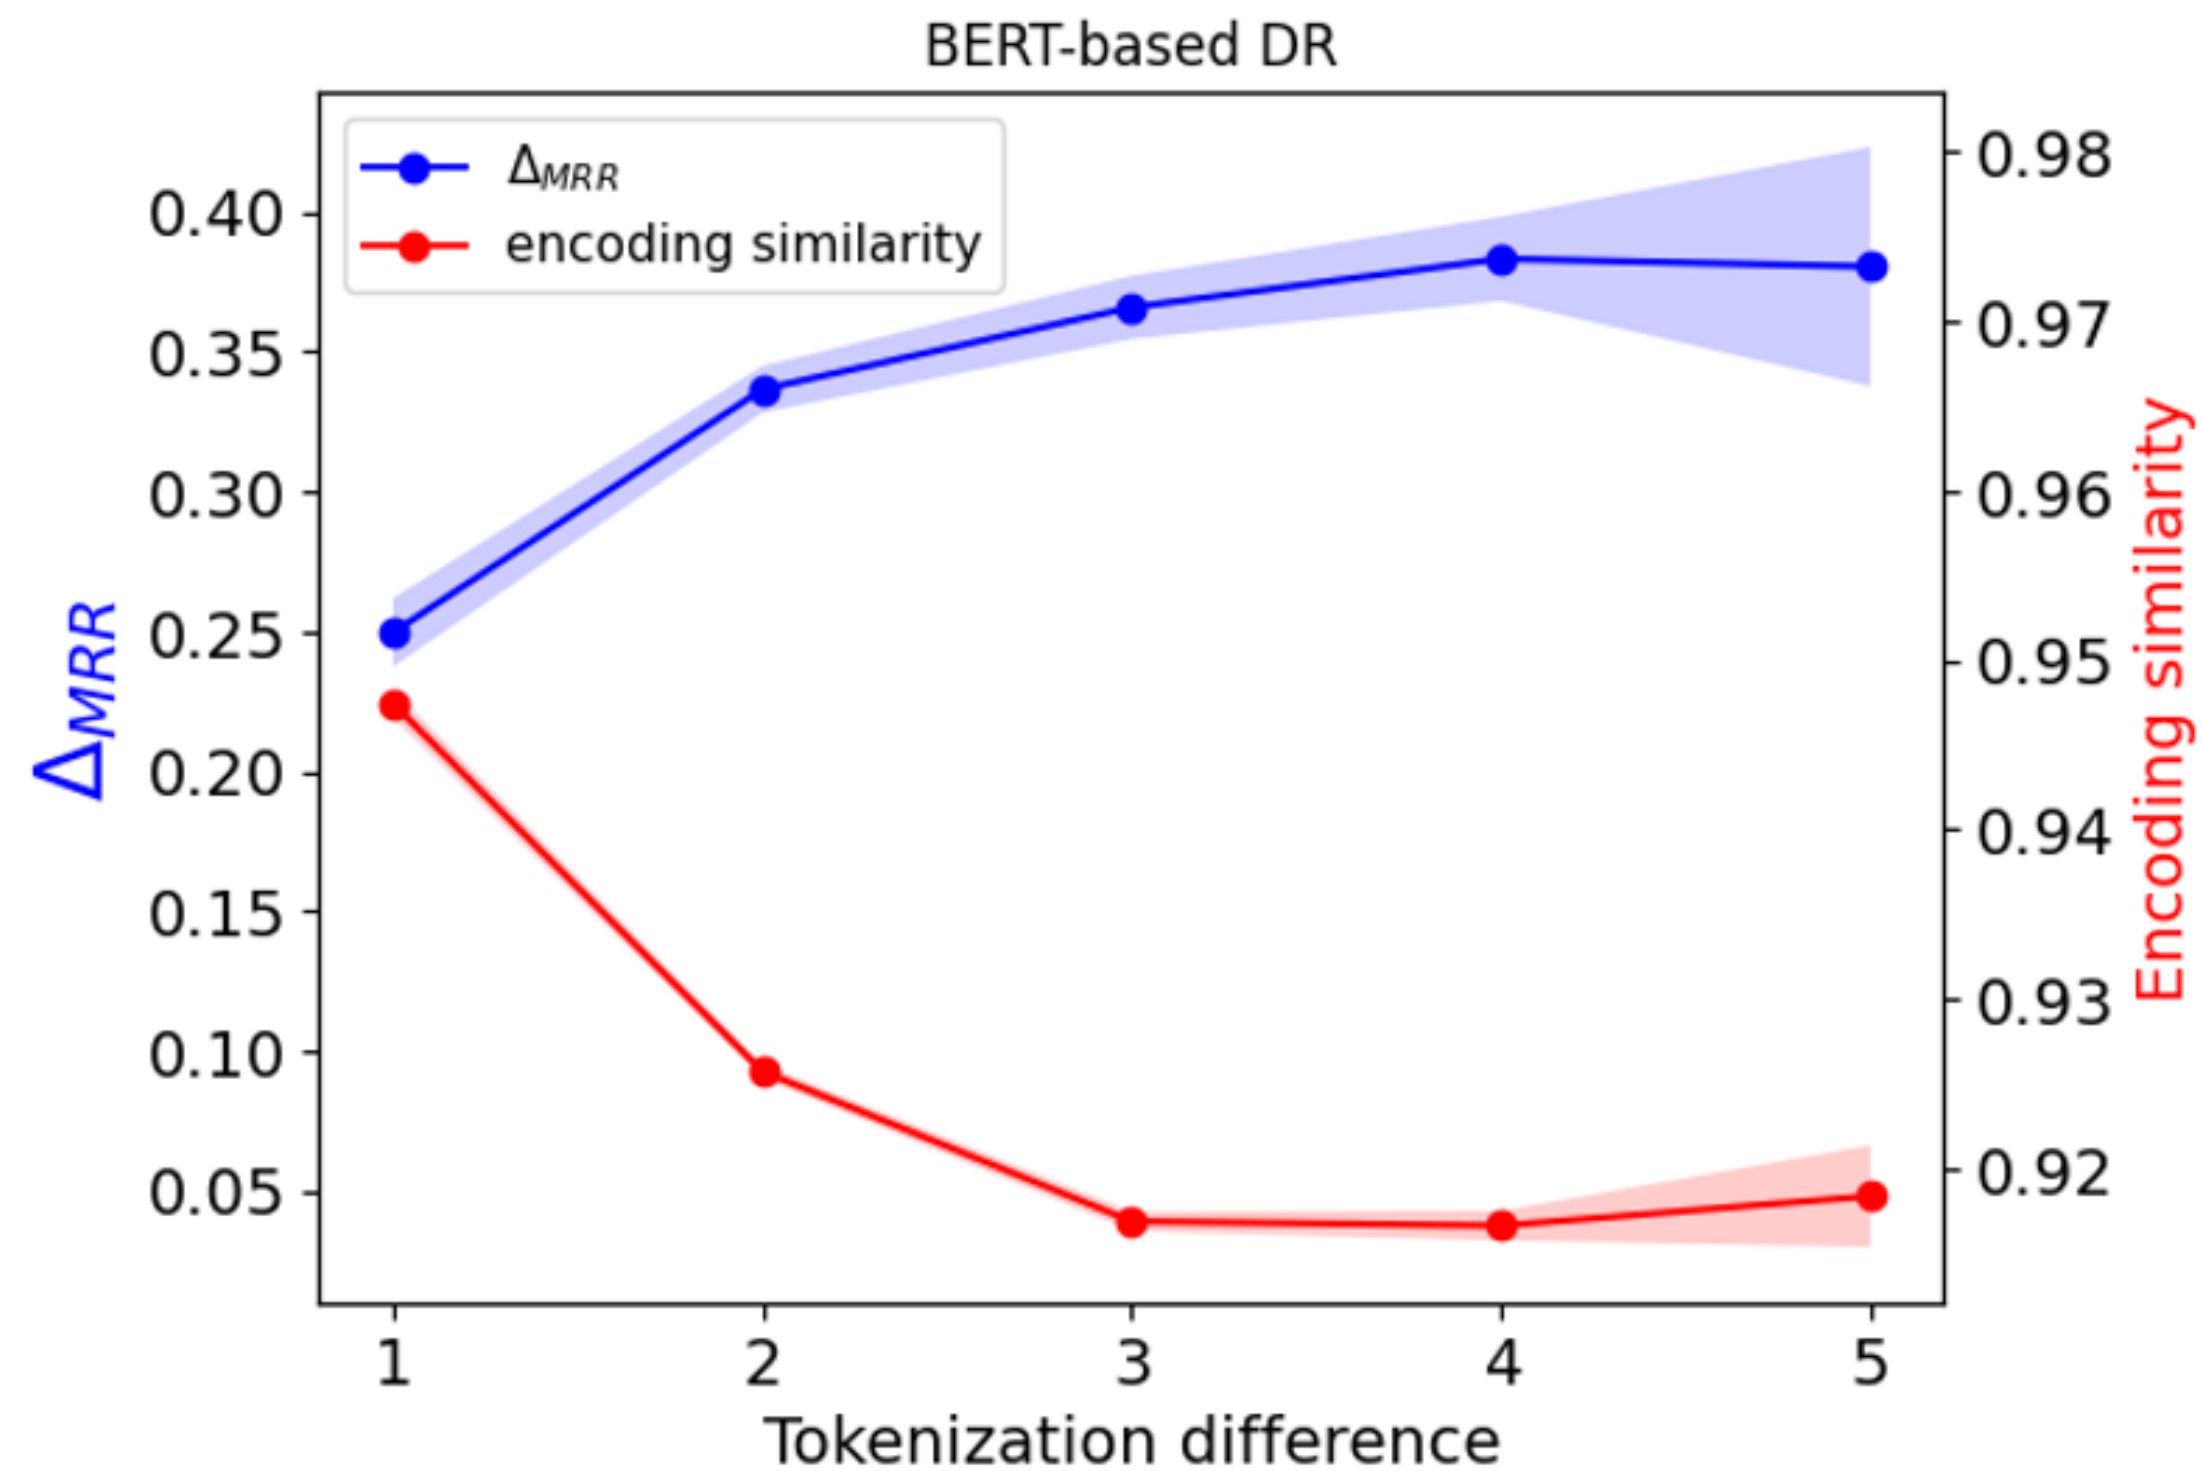


Typo query generation

- Synthetic Typo generation for MS MARCO queries
 - Random character Insertion: 'typo' -> 'tyapo'
 - Random character deletion: 'typo' -> 'tyo'
 - Random character substitution: 'typo' -> 'type'
 - Swap neighbour character: 'typo' -> 'tyop'
 - Swap adjacent keyboard character: 'typo' -> 'typi'
- These are common typos in real-world user queries [1]

[1] Matthias Hagen, Martin Potthast, Marcel Gohsen, Anja Rathgeber, and Benno Stein. 2017. A large-scale query spelling correction corpus. SIGIR

CharacterBERT-based DR behaves differently



Similar trends in other datasets

- We experimented with **MS MARCO** (dev queries), **TREC 2019** and **2020**
- We created a **new dataset**, DL-typo
 - 60 queries with typos from AOL query log
Human spelling corrections
 - Relevance assessments following TREC DL relevance criteria
On average 63.52 judgements per query (relevant: 25.7)